

# BoF: The IO-500 and the Virtual Institute of I/O

George Markomanolis

John Bent, Julian M. Kunkel, Jay Lofstead



# BoF Agenda

---

1. **BoF intro + The Virtual Institute for IO** (5 min) – Julian Kunkel
2. **What's new with IO-500** (8 min) – George Markomanolis
3. **Community lightning talks** (5 min each)
  - a. **In-node storage and memory-like I/O** — Adrian Jackson (EPCC)
  - b. **Demonstrating GPUDirect Storage using the IO500** — CJ Newburn/Sven Oehme
4. **Analysis of the IO-500 data** (12 min) – John Bent
5. **Award ceremony** (5 min) — George Markomanolis, John Bent, Julian Kunkel, Jay Lofstead
6. **Roadmap for the IO-500** (5 min) – Julian Kunkel
7. **Voice of the community & Open Discussion** (15 min) – Jay Lofstead

# What's new with IO-500

George S. Markomanolis,  
The IO-500 and the Virtual Institute of I/O  
Denver, Colorado, SC'19  
19 November 2019

ORNL is managed by UT-Battelle, LLC for the US Department of Energy



# Outline

- Benchmark improvements
  - MDTest Shift
  - Fix validation with IOR
- Versioning of IO500 via Tagging
  - How IO500 identifies proper versions of IOR/Mdtest
- Webpage: Info-Creator Drop-Down
- Student Cluster Competition 19

# MDtest Shift Added; IOR Shift Improved

- It was intentional to implement it since the beginning of IO-500
- Each process handles data that were accessed from another process
- It doesn't allow local caching
- It can hurt the performance significant
- In most cases, it is more realistic
- About IOR, improved shift how ranks mapped to nodes and IOR detects its mapping, no need to specify in IOR the mapping

# Versioning

- IOR/MDTest
  - We were using a tagging version lately
  - Now we use a HASH again inside prepare.sh
- IO500
  - We have a tag of the io-500-dev branch for sc19, ...
    - Tag always points to the latest version
    - Changelog shows details of changes
      - SC19-v1 and in Git commit message is the name as well
  - Trying to keep the instructions the same, at least per each list
  - Keep the versions of each IO500 submission through tagging
  - Needs to be improved though

# Info-Creator Drop-Down

- <https://www.vi4io.org/io500-info-creator/>

## Metadata server information

Number of nodes

Number of storage devices in each metadata server

Type of the storage media in metadata servers

-- select an option --

- SSD
- HDD
- NVMe
- NLSAS
- SAS
- SAS-SSD
- PD-SSD
- Other

Volatile memory capacity

Storage interface used by the servers

Network interconnect on the servers

File system software version on the servers

Operating system software version on the servers

Overhead of resilience in %

Type of the storage media in metadata servers

Other

Comment

## Data server information

Number of data server nodes

Number of storage devices

Type of the storage media

-- select an option --

Volatile memory capacity

Storage interface

-- select an option --

Network interconnect

File system software version

-- select an option --

Operating system software version

Overhead of resilience in %

Whatever

Comment

# Supercomputing Student Cluster Competition 2019

- IO-500 is part of the Supercomputing Cluster Competition 2019 for extra credits!
- New stonewall rule only for the competitions (30 seconds)
- Drop cache option for single node submissions
- We show that IO is important and should be considered part of such competitions
- New list for such competitions will be announced this week
- For vendors: If IO-500 becomes part of SCC, maybe you would like to provide hardware for a team

# The new IO-500 list and analysis

IO 500

# Reminder about Computing the Scores

- IOR easy
    - Write and read
  - IOR hard
    - Write and read
  - Mdttest easy
    - Create, stat, delete
  - Mdttest hard
    - Create, read, stat, delete
  - Namespace search
    - Find across all produced files
- 
- Geo mean -> Bandwidth score
- Geo mean -> Metadata score
- Geo mean -> Overall score

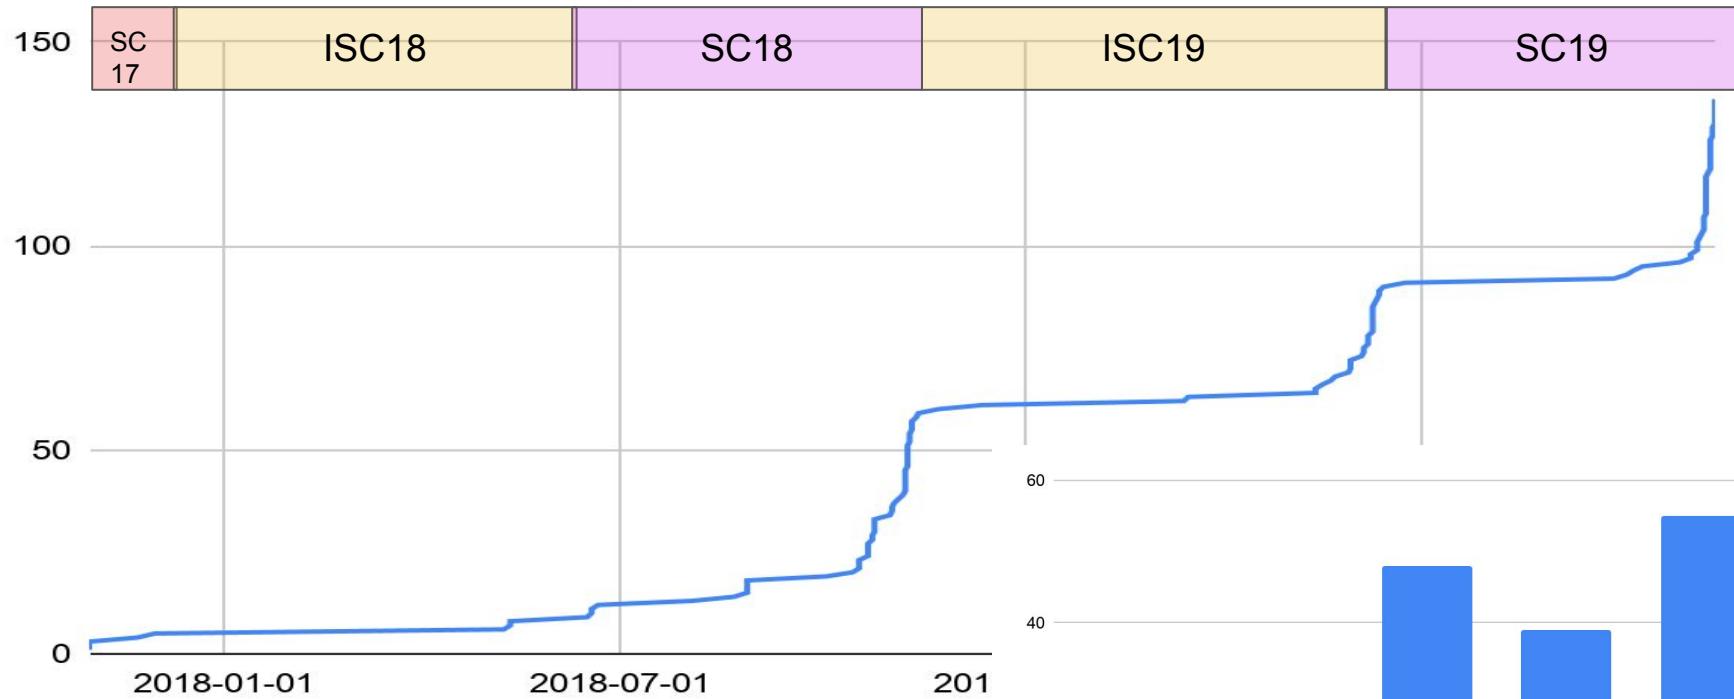
New metrics used for today's analysis:

`mdt_consume=geo_mean(mdt_easy_stat, mdt_easy_delete, mdt_hard_read, mdt_hard_stat, mdt_hard_delete)`

`mdt_produce=geo_mean(mdttest_easy_create, mdttest_hard_create)`

`mdt_ratio = mdt_consume / mdt_produce`

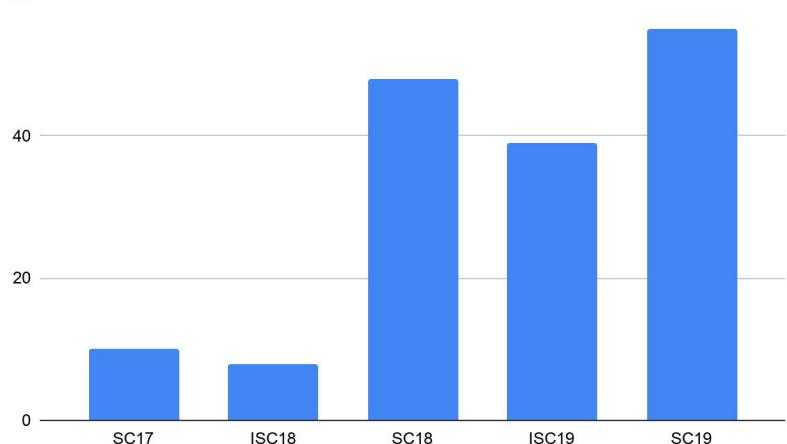
# Total Submissions by Date



Three new file systems in SC19!

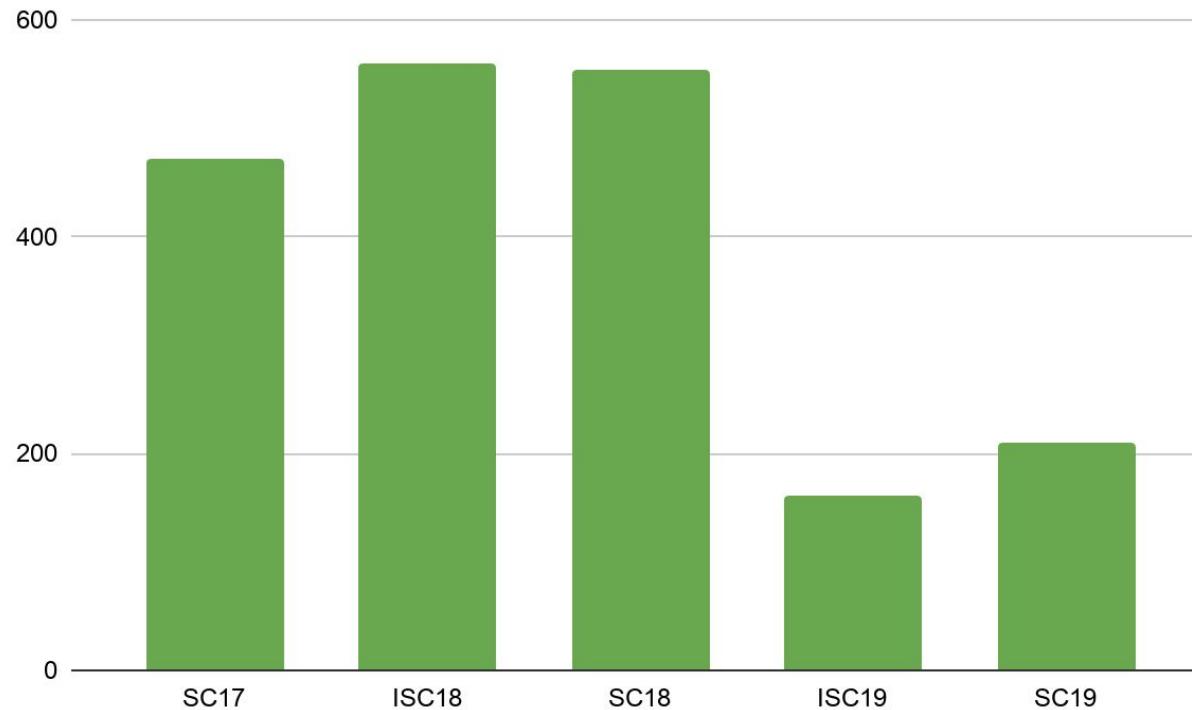
- DAOS
- GekkoFS
- YRCloudFile

SC19



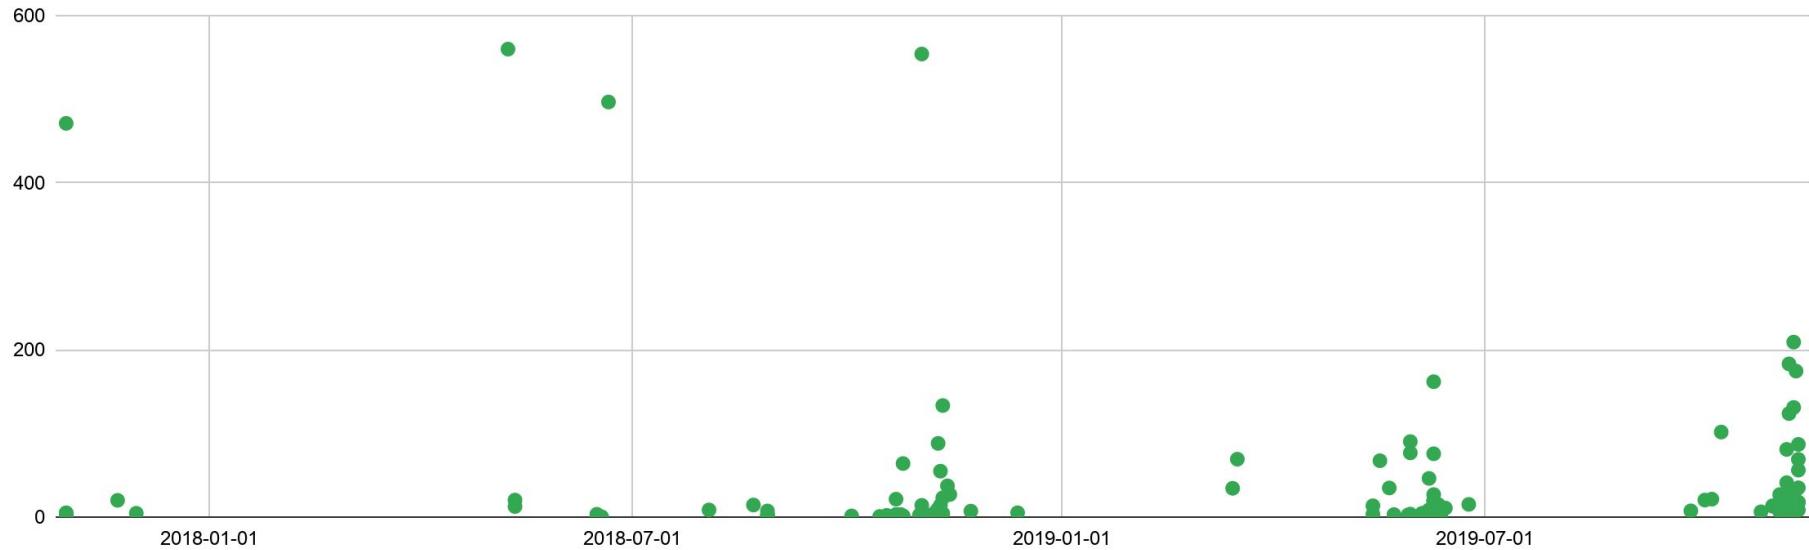
# Top Bandwidth by List

---



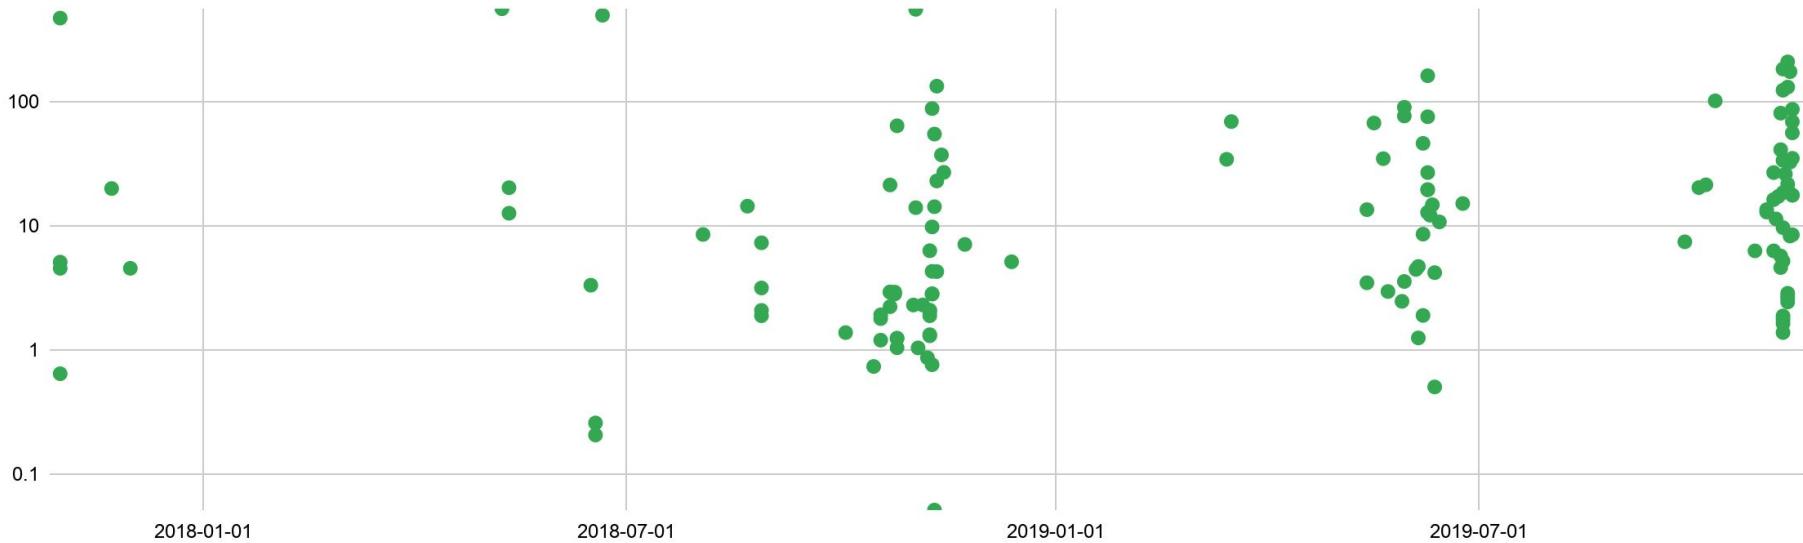
# All Bandwidths by Date

---



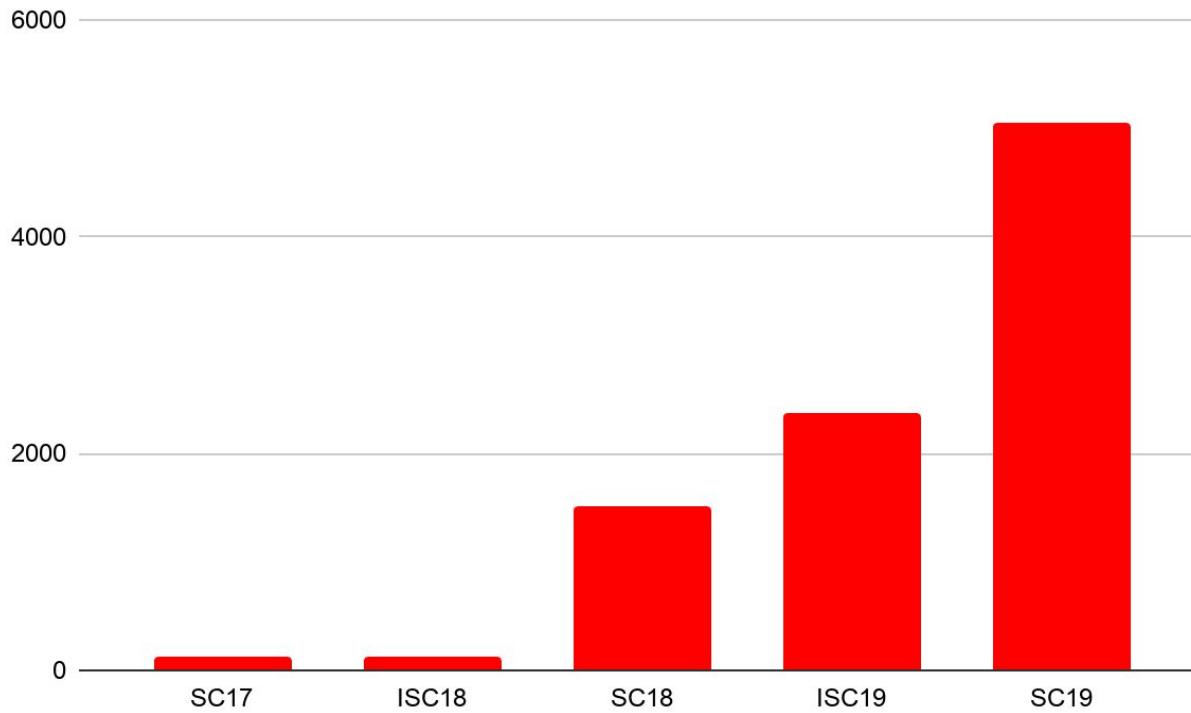
# All Bandwidths by Date (log-scale)

---



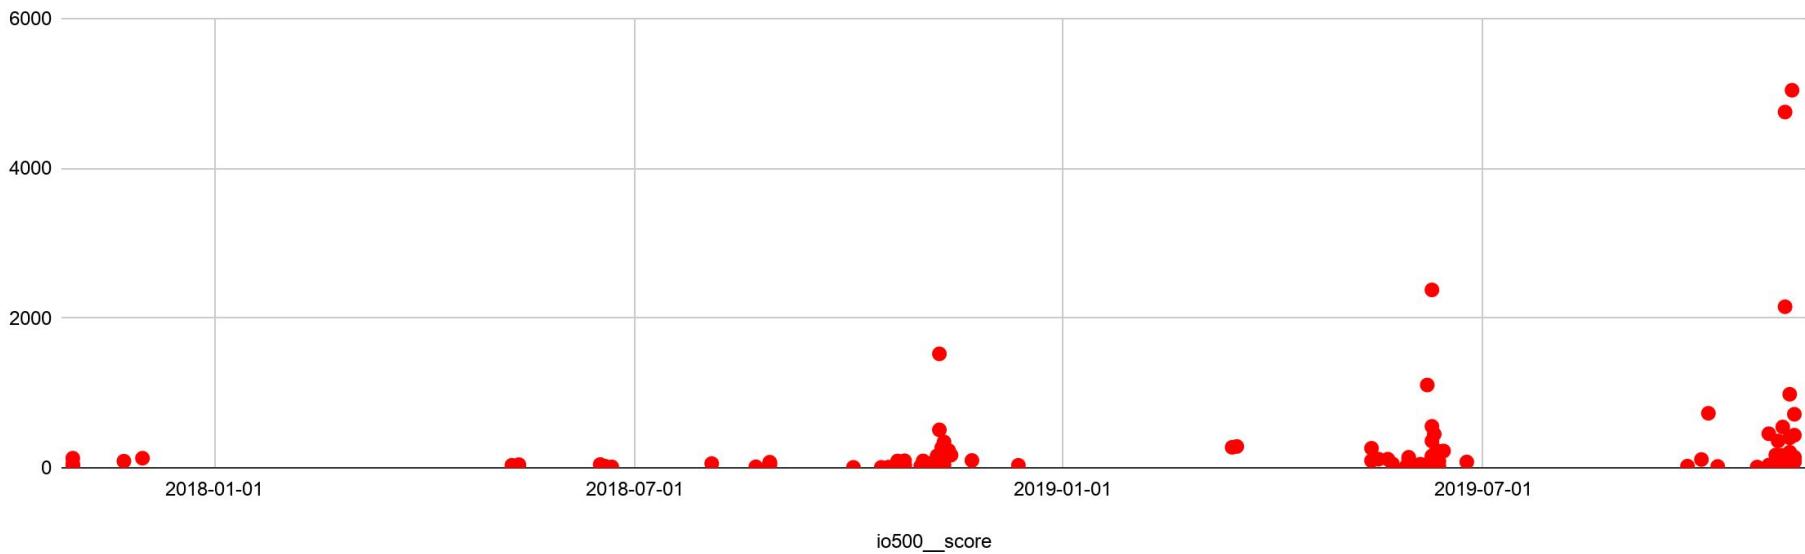
# Top Metadata by List

---



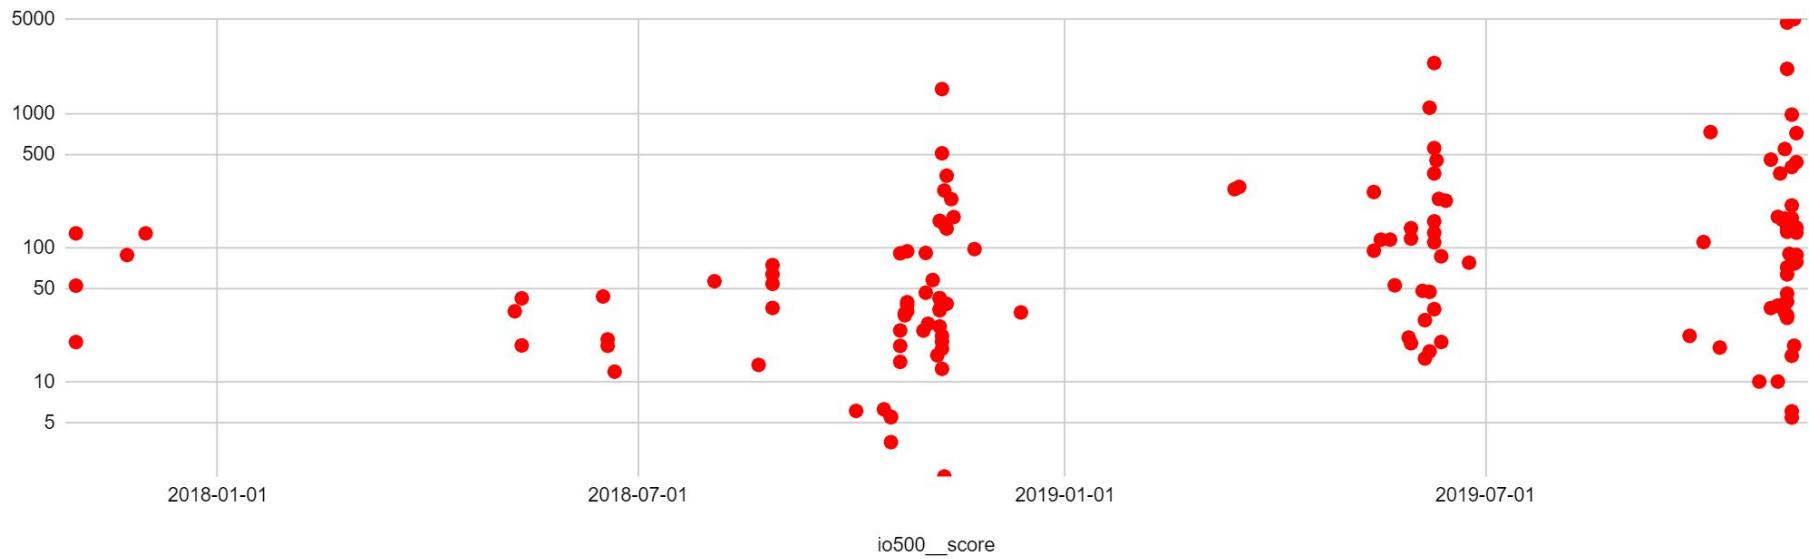
# All Metadata by Date

---



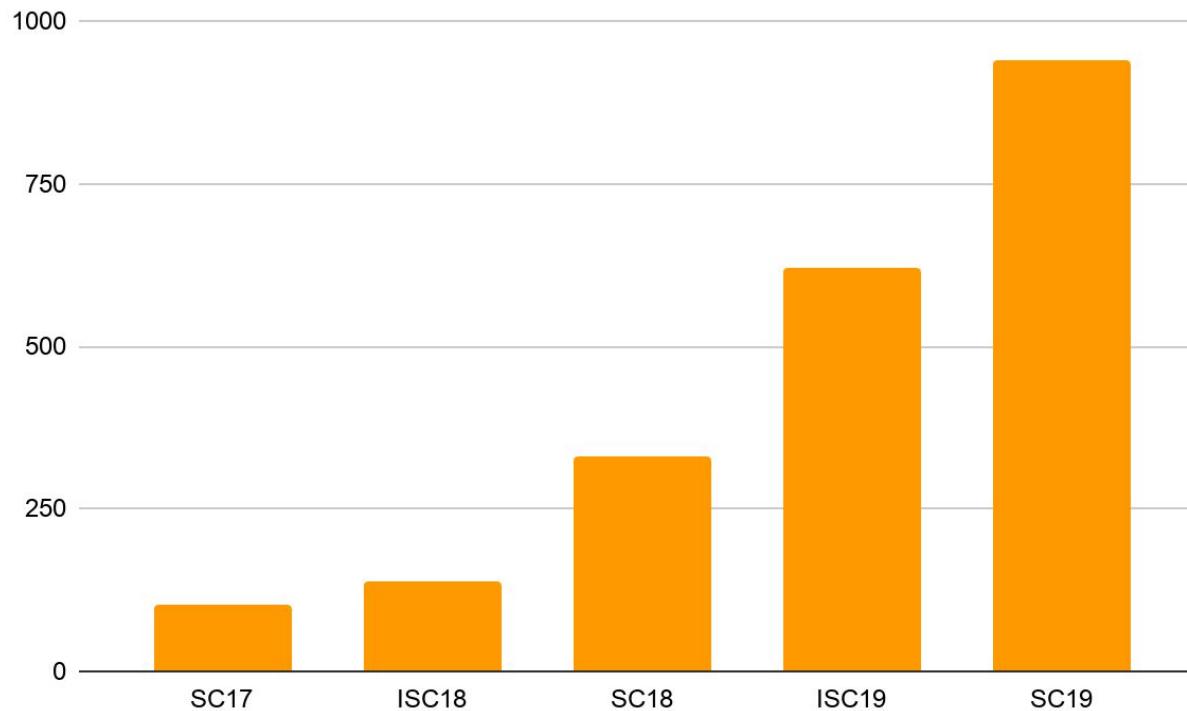
# All Metadata by Date (log-scale)

---



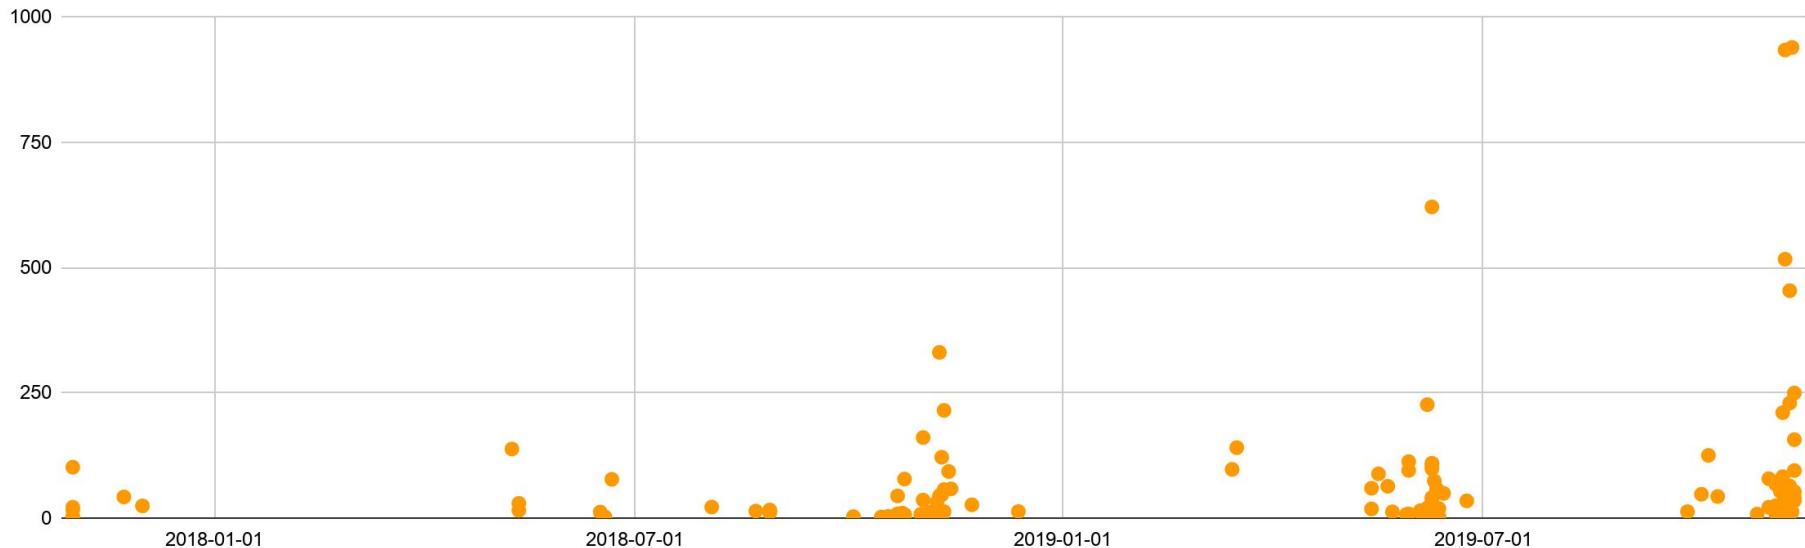
# Top Score by List

---



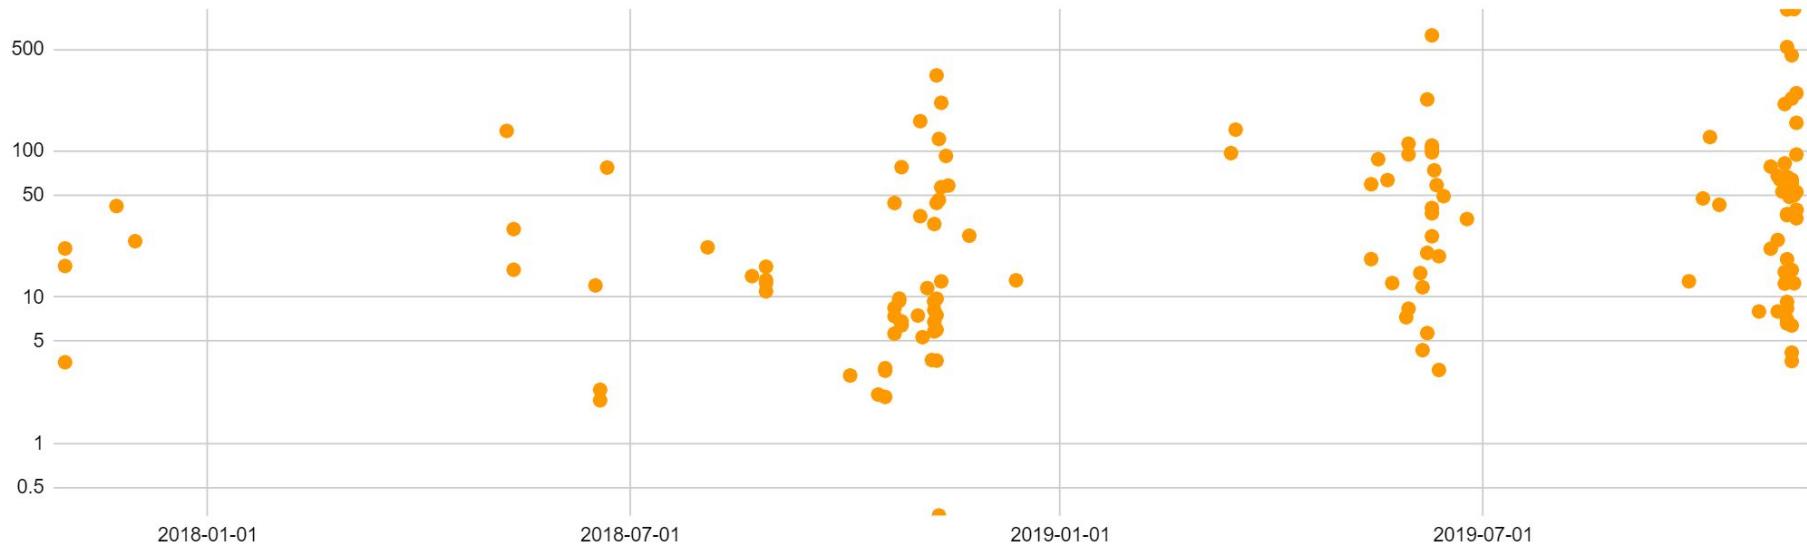
# All Scores by Date

---

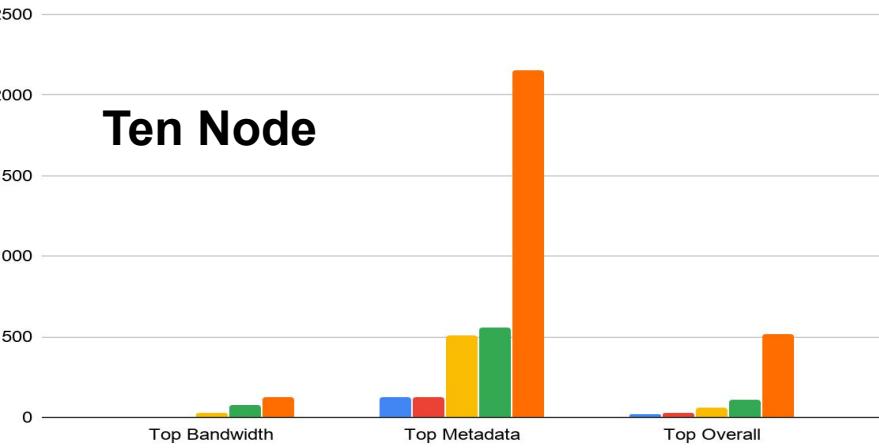


# All Scores by Date (log-scale)

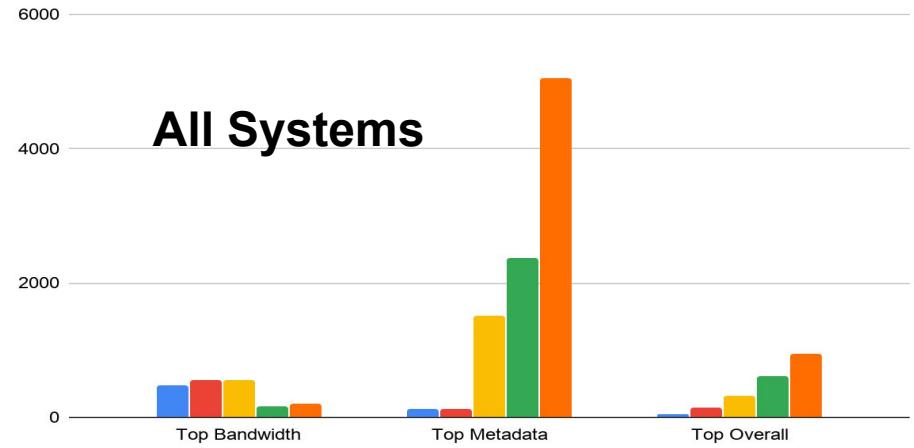
---



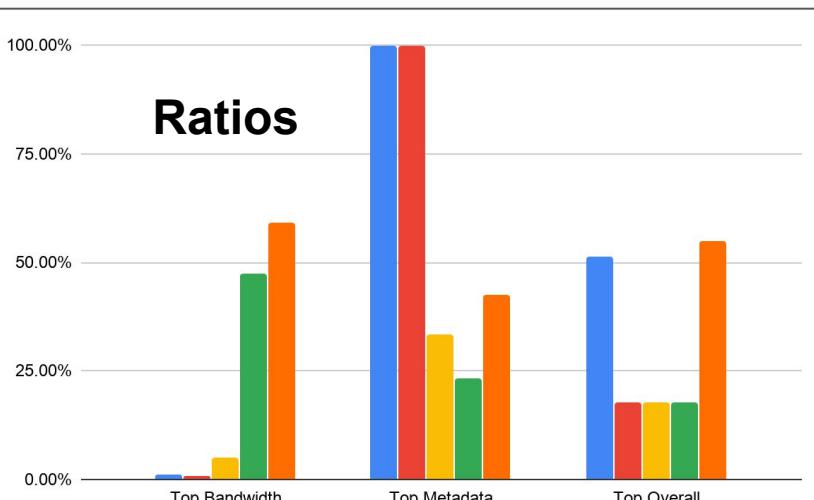
## Ten Node



## All Systems



## Ratios



## Top Scoring Systems by List

- SC17
- ISC18
- SC18
- ISC19
- SC19

# What to do with our lists? New Lists or Merge?

- Thus far, it seems like the new rules did not affect people's ability to improve scores
  - Suggests that perhaps we can just merge the new results into the old lists
  - But first let's consider a bit more carefully
- Let's zoom in on mdt\_produce, mdt\_consume, and mdt\_ratio
  - The new rule was designed out of fear that historical mdt\_consume rates artificially inflated by cache

Reminder:

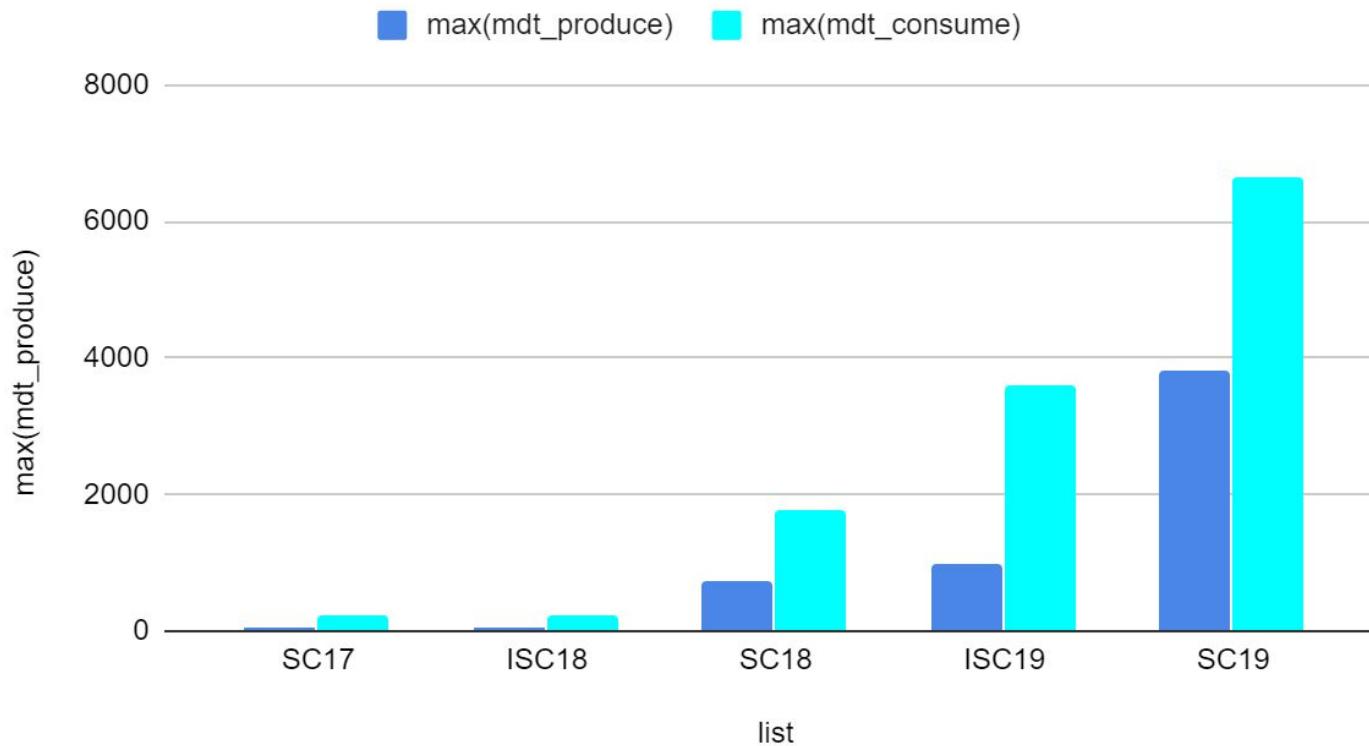
**mdt\_consume**=geo\_mean(mdt\_easy\_stat, mdt\_easy\_delete, mdt\_hard\_read, mdt\_hard\_stat, mdt\_hard\_delete)

**mdt\_produce**=geo\_mean(mdtest\_easy\_create, mdtest\_hard\_create)

**mdt\_ratio** = mdt\_consume / mdt\_produce

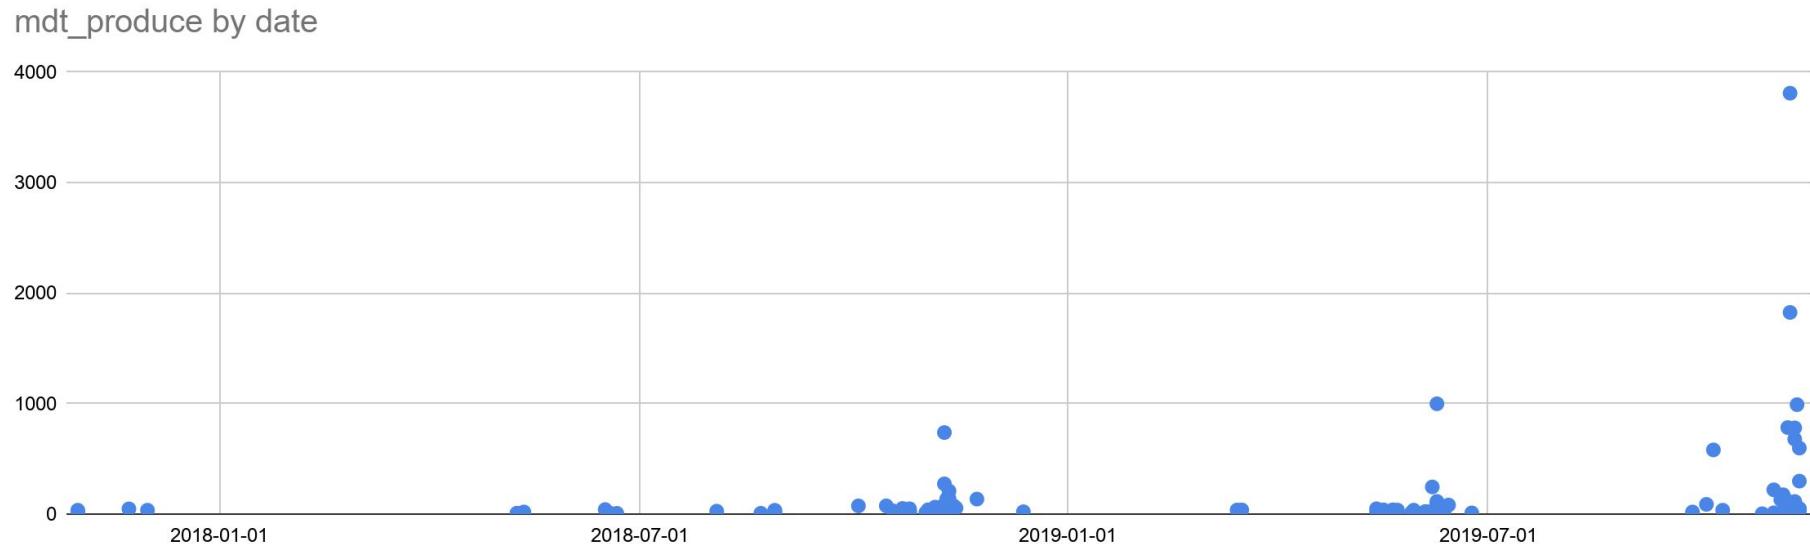
# Top mdt\_produce and mdt\_consume by List

---



# All mdt\_produce by Date

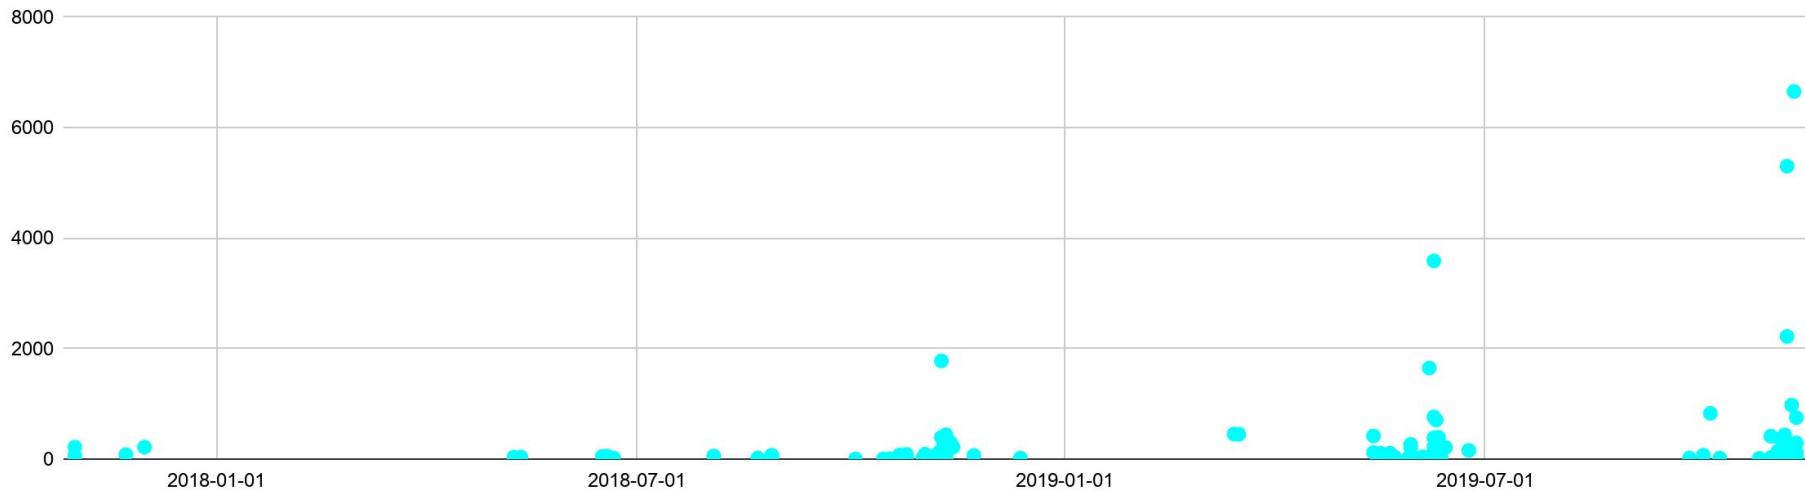
---



# All mdt\_consume by Date

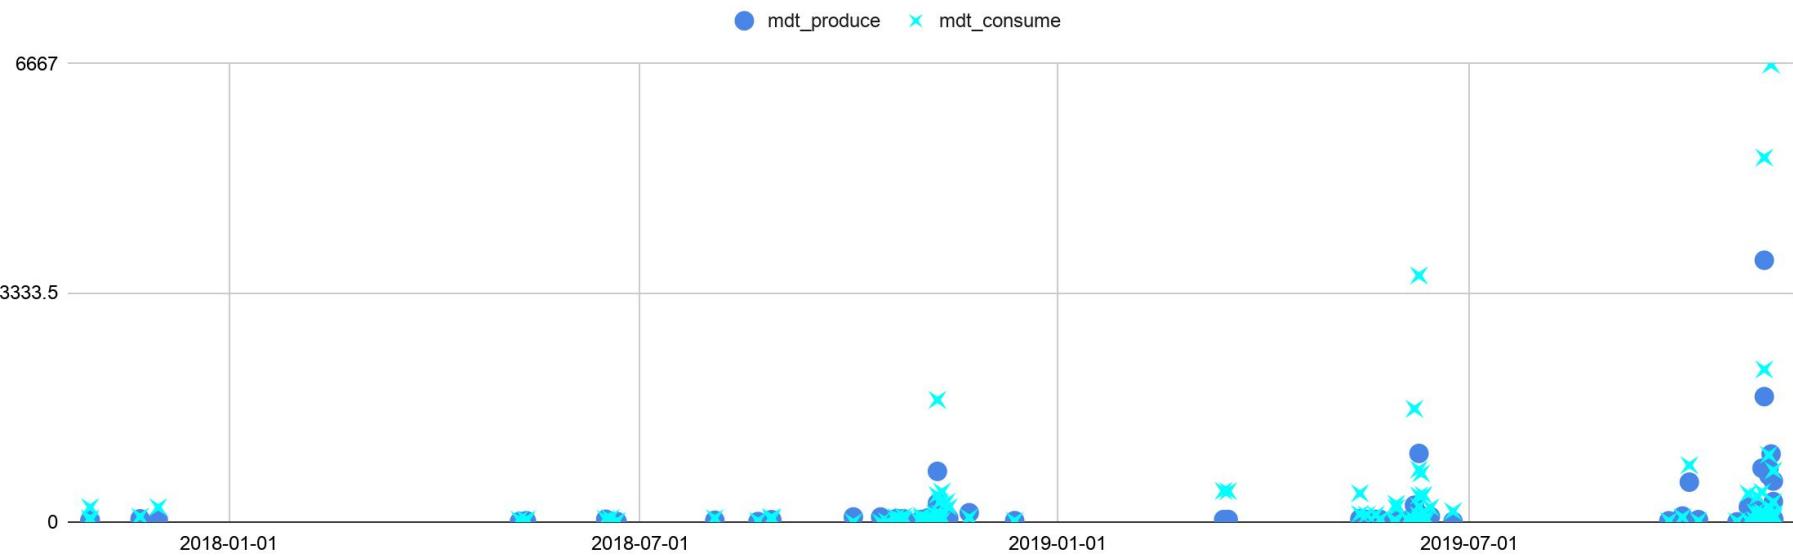
---

mdt\_consume by date



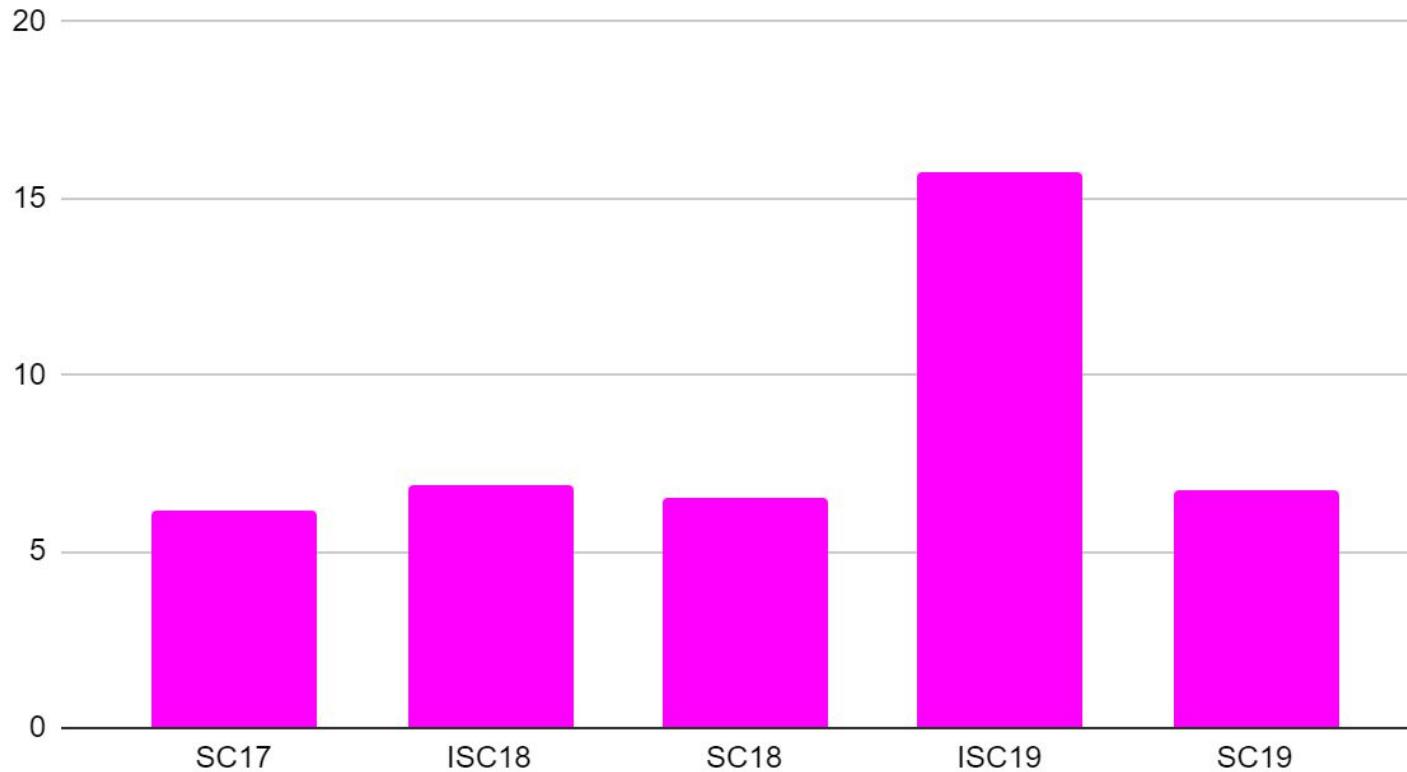
# All mdt\_produce and mdt\_consume by Date

---

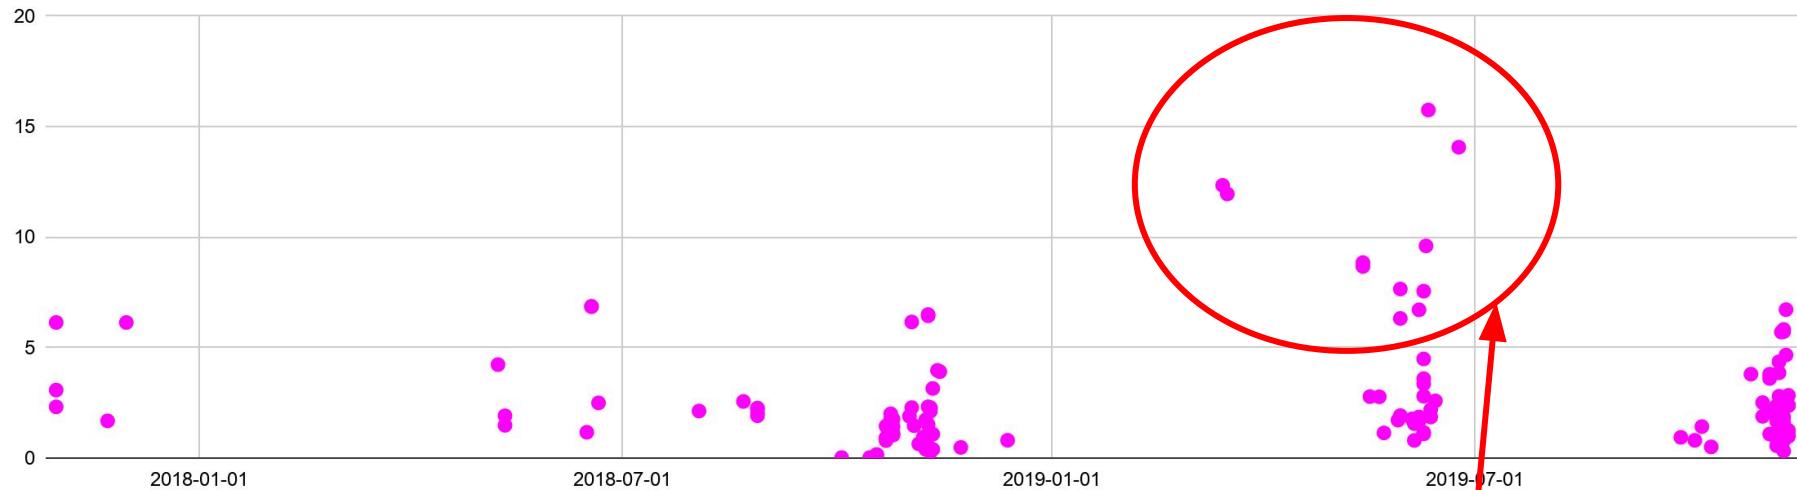


# Top mdt\_ratio by List

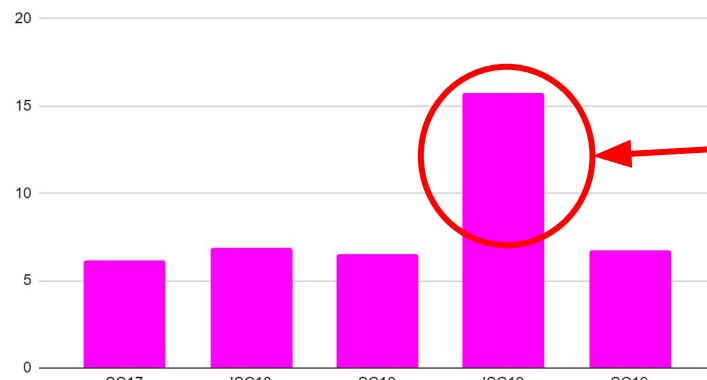
---



## Metadata Consume / Metadata Produce (mdt\_ratio) By Date



### Top Mdtest Consume / Mdtest Produce (mdt\_ratio) by List



- Something definitely changed in ISC19.  
Ratio is probably not expected to increase.

# Apples:Apples Comparison of Shift Effects

---

- Five people submitted apples to apples results
  - Ken Carlile, Andreas Dilger, Sven Oehme, Matt Rásó-Barnett, Ruibo Wang
  - Unfortunately four were for Lustre systems, Ken's was Vast however
- Reminder that shift was to avoid client side cache
  - IOR phases and mdtest produce phases and find should not be degraded
  - mdtest stat and read phases could be degraded since client side cache could help these
  - mdtest delete phase might be a bit less likely to be degraded since server must be involved

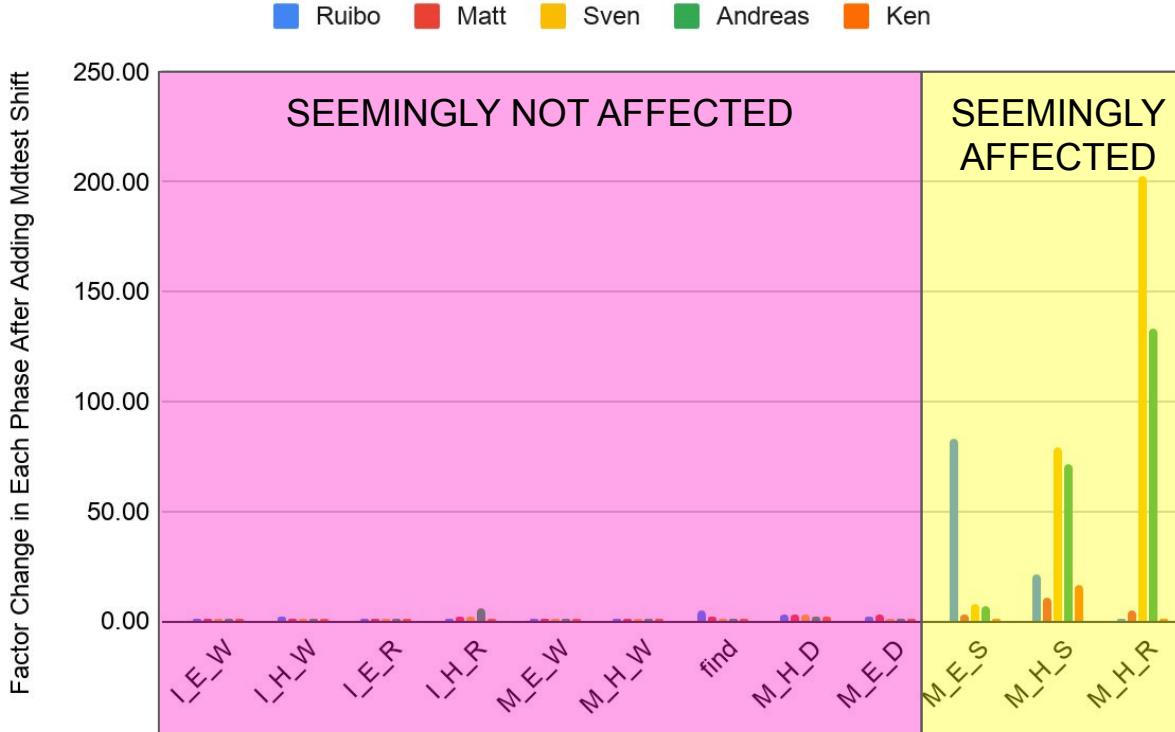
# Apples:Apples Comparison of mdtest Shift Effects

- Five people submitted apples to apples results
  - Ken Carlile, Andreas Dilger, Sven Oehme, Matt Rásó-Barnett, Ruibo Wang
- Reminder that shift was added to avoid client side cache
  - IOR phases and mdtest produce phases and find should not be degraded
  - mdtest stat and read phases could be degraded since client side cache could help these
  - mdtest delete phase might be a bit less likely to be degraded since server must be involved
  - If client-side caching had been helping, we would expect to see this in mdt\_ratio



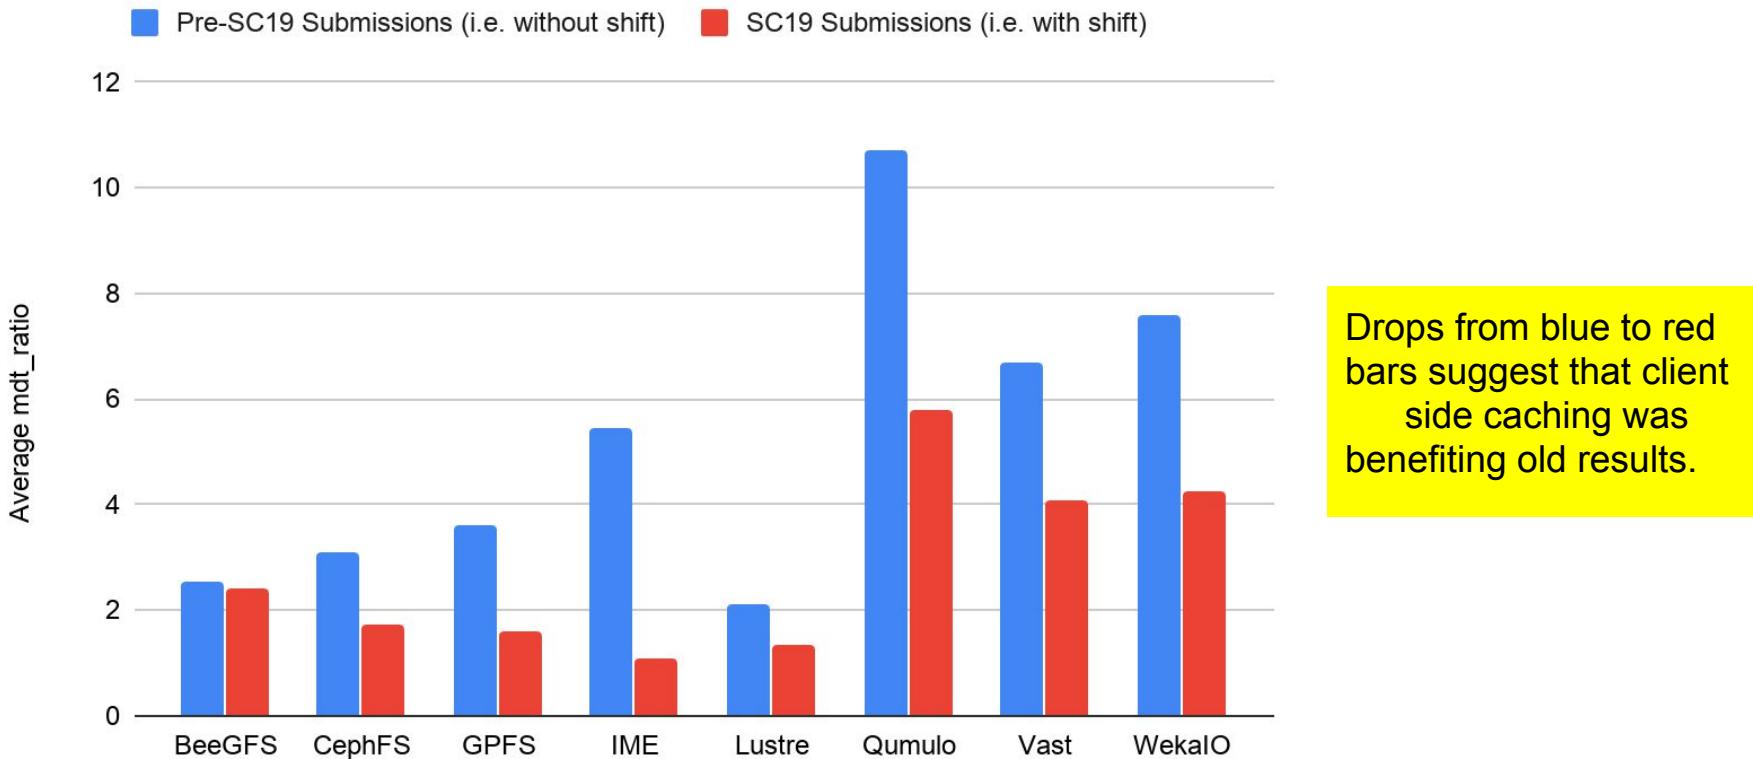
# Apples:Apples Results Show Effect of Change

I=IOR  
M=mdtest  
E=easy  
H=hard  
W=write  
R=read  
S=stat  
D=delete



But this was  
only two file  
systems!  
Ken was  
Vast; others  
were Lustre.

# 8 File Systems Had Both Historical and SC19 Results



# Clearly mdtest Shift Had an Effect - Will Create New Lists

---

Four Lists Will be Maintained Going Forward

1. Full List
  - a. No historical results
  - b. All submissions to SC19 and beyond
2. Ranked List
  - a. No historical results
  - b. Multiple submissions to SC19 and beyond for “system/institution/file system” collapsed into top
3. Ten Node Ranked List
  - a. No historical results
  - b. Only 10 node submissions to SC19 and beyond will be included
  - c. Multiple submissions for “system/institution/file system” collapsed into top submission
4. Historical List
  - a. All submissions both historical and new are included

Thanks to community members Ken Carlile, Andreas Dilger, Glenn Lockwood, Sven Oehme, Matt Rásó-Barnett, and Ruibo Wang for offering valuable opinions and data to help with this key decision!

# Awards

10<sup>500</sup>

# Six SC19 Awards Will Be Now Given

---

1. Ten-node
  - a. Bandwidth
  - b. Metadata
  - c. Overall
2. All Systems
  - a. Bandwidth
  - b. Metadata
  - c. Overall

Note that due to the decision about making a new list, only SC19 submissions can compete.

Even though historical bandwidths are fully compatible and were not affected by the mdtest-shift, the committee decided that the move to a new list should be complete. This will minimize any confusion as well as reduce the likelihood that incompatible results are ever inadvertently compared.

# 10 node challenge - Bandwidth Winner

## 10 Node Challenge SC19 ONLY

This is the official list from  Supercomputing 2019 for the 10 Node Challenge. The list shows the best result for a given combination of system/institution/filesystem qualifying for the 10 Node Challenge.

IO 500

Sorted by BW

#	information								io500	bw	GiB/s
	list id	institution	system	storage vendor	filesystem type	client nodes	client total procs	data			
1	sc19	Intel	Wolf	Intel	DAOS	10	310	zip		123.89	
2	sc19	National Supercomputing Centre, Singapore	Aspire 1	DDN	IME	10	160	zip		101.75	
3	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	zip		86.97	
4	sc19	WekaIO	WekaIO	WekaIO	WekaIO Matrix	10	2610	zip		56.22	
5	sc19	State Key Laboratory of High-end Server & Storage Technology (HSS)	TStor3000	INSPUR	BeeGFS	10	300	zip		41.14	
6	sc19	National Supercomputing Center in Changsha	Tianhe-2E	National University of Defense Technology	Lustre	10	160	zip		35.06	
7	sc19	CSIRO	bracewell	Dell/ThinkParQ	beegfs	10	160	zip		33.77	
8	sc19	Janelia Research Campus, HHMI	weka	WekaIO	wekaio	18	1368	zip		26.22	
9	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	10	320	zip		21.73	
10	sc19	EPCC	NEXTGenIO Prototype	BSC (NEXTGenIO) & JGU (Ada-FS)	Adhoc Filesystem	10	240	zip		21.47	

IO  
IO

IO<sup>500</sup>

# Certificate

## IO-500 Performance Certification

This Certificate is awarded to:

**Intel**

#1 in the 10 Node BW Score

IO<sup>500</sup>



Nov 2019

IO-500 Steering Board

<http://io500.org/list/19-11/>

IO  
IO

# 10 node challenge - Metadata Winner

## 10 Node Challenge SC19 ONLY

This is the official list from  Supercomputing 2019 for the 10 Node Challenge. The list shows the best result for a given combination of system/institution/filesystem qualifying for the 10 Node Challenge.



Sorted by md

#	list id	information									io500	
		institution	system	storage vendor	filesystem type	client nodes	client total procs	data	bw	md	GiB/s	kIOP/s
									bw	md		
1	sc19	Intel	Wolf	Intel	DAOS	10	310	zip	123.89	2152.46		
2	sc19	EPCC	NEXTGenIO Prototype	BSC (NEXTGenIO) & JGU (Ada-FS)	Adhoc Filesystem	10	240	zip	21.47	728.68		
3	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	zip	86.97	715.76		
4	sc19	iFLYTEK	iFLYTEK	Yanrong	YRCloudFile	10	200	zip	13.55	455.18		
5	sc19	WekalO	WekalO	WekalO	WekalO Matrix	10	2610	zip	56.22	435.76		
6	sc19	DDN	AI400	DDN	Lustre	10	240	zip	19.65	207.63		
7	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	10	320	zip	21.73	167.09		
8	sc19	State Key Laboratory of High-end Server & Storage Technology (HSS)	TStor3000	INSPUR	BeeGFS	10	300	zip	41.14	165.71		
9	sc19	CSIRO	bracewell	Dell/ThinkParQ	beegfs	10	160	zip	33.77	132.15		
10	sc19	Janelia Research Campus, HHMI	weka	WekalO	wekaio	18	1368	zip	26.22	90.62		

IO  
IO

IO<sup>500</sup>

# Certificate

## IO-500 Performance Certification

This Certificate is awarded to:

**Intel**

#1 in the 10 Node MD Score

IO<sup>500</sup>



Nov 2019

10-500 Steering Board

<http://io500.org/list/19-11/>

IO  
IO

# 10 node challenge - Winner

## 10 Node Challenge SC19 ONLY

This is the official list from  Supercomputing 2019 for the 10 Node Challenge. The list shows the best result for a given combination of system/institution/filesystem qualifying for the 10 Node Challenge.

IO 500

Sorted by score

#	information								io500		
	list id	institution	system	storage vendor	filesystem type	client nodes	client total procs	data	score	bw	md
									GiB/s	kIOP/s	
1	sc19	Intel	Wolf	Intel	DAOS	10	310	zip	516.41	123.89	2152.46
2	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	zip	249.50	86.97	715.76
3	sc19	WekaIO	WekaIO	WekaIO	WekaIO Matrix	10	2610	zip	156.51	56.22	435.76
4	sc19	EPCC	NEXTGenIO Prototype	BSC (NEXTGenIO) & JGU (Ada-FS)	Adhoc Filesystem	10	240	zip	125.08	21.47	728.68
5	sc19	State Key Laboratory of High-end Server & Storage Technology (HSS)	TStor3000	INSPUR	BeeGFS	10	300	zip	82.57	41.14	165.71
6	sc19	iFLYTEK	iFLYTEK	Yanrong	YRCloudFile	10	200	zip	78.54	13.55	455.18
7	sc19	CSIRO	bracewell	Dell/ThinkParQ	beegfs	10	160	zip	66.80	33.77	132.15
8	sc19	DDN	AI400	DDN	Lustre	10	240	zip	63.88	19.65	207.63
9	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	10	320	zip	60.25	21.73	167.09
10	sc19	National Supercomputing Center in Changsha	Tianhe-2E	National University of Defense Technology	Lustre	10	160	zip	52.58	35.06	78.86

IO  
IO

IO<sup>500</sup>

# Certificate

## IO-500 Performance Certification

This Certificate is awarded to:

**Intel**

#1 in the 10 Node Challenge

IO<sup>500</sup>



Nov 2019

IO-500 Steering Board

<http://io500.org/list/19-11/>

IO  
IO

# Full list - Bandwidth Winner

## Full List

This is the full list from Supercomputing 2019. The list shows all submissions.

**IO 500**  
Sorted by BW

#	information								io500	
	list id	institution	system	storage vendor	filesystem type	client nodes	client total procs	data	bw	GiB/s
1	sc19	National Supercomputing Center in Changsha	Tianhe-2E	National University of Defense Technology	Lustre	480	5280	zip	209.43	
2	sc19	Intel	Wolf	Intel	DAOS	26	728	zip	183.36	
3	sc19	WekaIO	WekaIO	WekaIO	WekaIO Matrix	345	8625	zip	174.74	
4	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	128	2048	zip	131.25	
5	sc19	Intel	Wolf	Intel	DAOS	10	310	zip	123.89	
6	sc19	National Supercomputing Centre, Singapore	Aspire 1	DDN	IME	10	160	zip	101.75	
7	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	zip	86.97	
8	sc19	CEA	Tera-1000	DDN	Lustre	128	4096	zip	81.01	
9	sc19	CSIRO	bracewell scratch2	Dell/ThinkParQ	beegfs	26	260	zip	69.10	
10	sc19	WekaIO	WekaIO	WekaIO	WekaIO Matrix	10	2610	zip	56.22	

IO  
IO

IO<sup>500</sup>

# Certificate

IO-500 Performance Certification

This Certificate is awarded to:

**National Supercomputing Center in Changsha**

#1 in the IO-500 BW Score

IO 500



Nov 2019

IO-500 Steering Board

<http://io500.org/list/19-11/>

IO  
IO

# Full list - Metadata Winner

## Full List

This is the full list from  Supercomputing 2019. The list shows all submissions.

IO 500

Sorted by md

#	information								io500	
	list id	institution	system	storage vendor	filesystem type	client nodes	client total procs	data	bw	md
									GiB/s	kIOP/s
1	sc19	WekalO	WekalO	WekalO	WekalO Matrix	345	8625	<a href="#">zip</a>	174.74	5045.33
2	sc19	Intel	Wolf	Intel	DAOS	26	728	<a href="#">zip</a>	183.36	4753.79
3	sc19	Intel	Wolf	Intel	DAOS	10	310	<a href="#">zip</a>	123.89	2152.46
4	sc19	National Supercomputing Center in Changsha	Tianhe-2E	National University of Defense Technology	Lustre	480	5280	<a href="#">zip</a>	209.43	982.78
5	sc19	EPCC	NEXTGenIO Prototype	BSC (NEXTGenIO) & JGU (Ada-FS)	Adhoc Filesystem	10	240	<a href="#">zip</a>	21.47	728.68
6	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	<a href="#">zip</a>	86.97	715.76
7	sc19	CEA	Tera-1000	DDN	Lustre	128	4096	<a href="#">zip</a>	81.01	545.74
8	sc19	iFLYTEK	iFLYTEK	Yanrong	YRCloudFile	10	200	<a href="#">zip</a>	13.55	455.18
9	sc19	WekalO	WekalO	WekalO	WekalO Matrix	10	2610	<a href="#">zip</a>	56.22	435.76
10	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	128	2048	<a href="#">zip</a>	131.25	401.13

IO  
IO

IO<sup>500</sup>

# Certificate

## IO-500 Performance Certification

This Certificate is awarded to:

# WekaIO

#1 in the IO-500 MD Score

IO 500



Nov 2019

IO-500 Steering Board

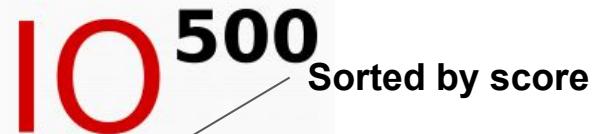
<http://io500.org/list/19-11/>

IO  
IO

# Full list - Winner

## Full List

This is the full list from  Supercomputing 2019. The list shows all submissions.



#	information								io500		
	list id	institution	system	storage vendor	filesystem type	client nodes	client total procs	data	score	bw	md
									GiB/s	kIOP/s	
1	sc19	WekalO	WekalO	WekalO	WekalO Matrix	345	8625	zip	938.95	174.74	5045.33
2	sc19	Intel	Wolf	Intel	DAOS	26	728	zip	933.64	183.36	4753.79
3	sc19	Intel	Wolf	Intel	DAOS	10	310	zip	516.41	123.89	2152.46
4	sc19	National Supercomputing Center in Changsha	Tianhe-2E	National University of Defense Technology	Lustre	480	5280	zip	453.68	209.43	982.78
5	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	zip	249.50	86.97	715.76
6	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	128	2048	zip	229.45	131.25	401.13
7	sc19	CEA	Tera-1000	DDN	Lustre	128	4096	zip	210.26	81.01	545.74
8	sc19	WekalO	WekalO	WekalO	WekalO Matrix	10	2610	zip	156.51	56.22	435.76
9	sc19	EPCC	NEXTGenIO Prototype	BSC (NEXTGenIO) & JGU (Ada-FS)	Adhoc Filesystem	10	240	zip	125.08	21.47	728.68
10	sc19	CSIRO	bracewell scratch2	Dell/ThinkParQ	beegfs	26	260	zip	94.86	69.10	130.23

# Full list - Winner

## Full List

This is the full list from Supercomputing 2019. The list shows all submissions.

**IO 500** Sorted by score

#	information								io500		
	list id	institution	system	storage vendor	filesystem type	client nodes	client total procs	data	score	bw	md
									GiB/s	kIOP/s	
1	sc19	WekalO	WekalO	WekalO	WekalO Matrix	345	8625	zip	938.95	174.74	5045.33
2	sc19	Intel	Wolf	Intel	DAOS	26	728	zip	933.64	183.36	4753.79
3	sc19	Intel	Wolf	Intel	DAOS	10	310	zip	516.41	123.89	2152.46
4	sc19	National Supercomputing Center in Changsha	Tianhe-2E	National University of Defense Technology	Lustre	480	5280	zip	453.68	209.43	982.78
5	sc19	NVIDIA	DGX-2H SuperPOD	DDN	Lustre	10	400	zip	249.50	86.97	715.76
6	sc19	University of Cambridge	Data Accelerator	Dell EMC	Lustre	128	2048	zip	229.45	131.25	401.13
7	sc19	CEA	Tera-1000	DDN	Lustre	128	4096	zip	210.26	81.01	545.74
8	sc19	WekalO	WekalO	WekalO	WekalO Matrix	10	2610	zip	156.51	56.22	435.76
9	sc19	EPCC	NEXTGenIO Prototype	BSC (NEXTGenIO) & JGU (Ada-FS)	Adhoc Filesystem	10	240	zip	125.08	21.47	728.68
10	sc19	CSIRO	bracewell scratch2	Dell/ThinkParQ	beegfs	26	260	zip	94.86	69.10	130.23

0.57% difference

IO IO IO IO IO IO IO IO IO

# Certificate

## IO-500 Performance Certification

## This Certificate is awarded to:

# WekaIO

#1 in the IO-500

IO 500



Nov 2019

## 10-500 steering board

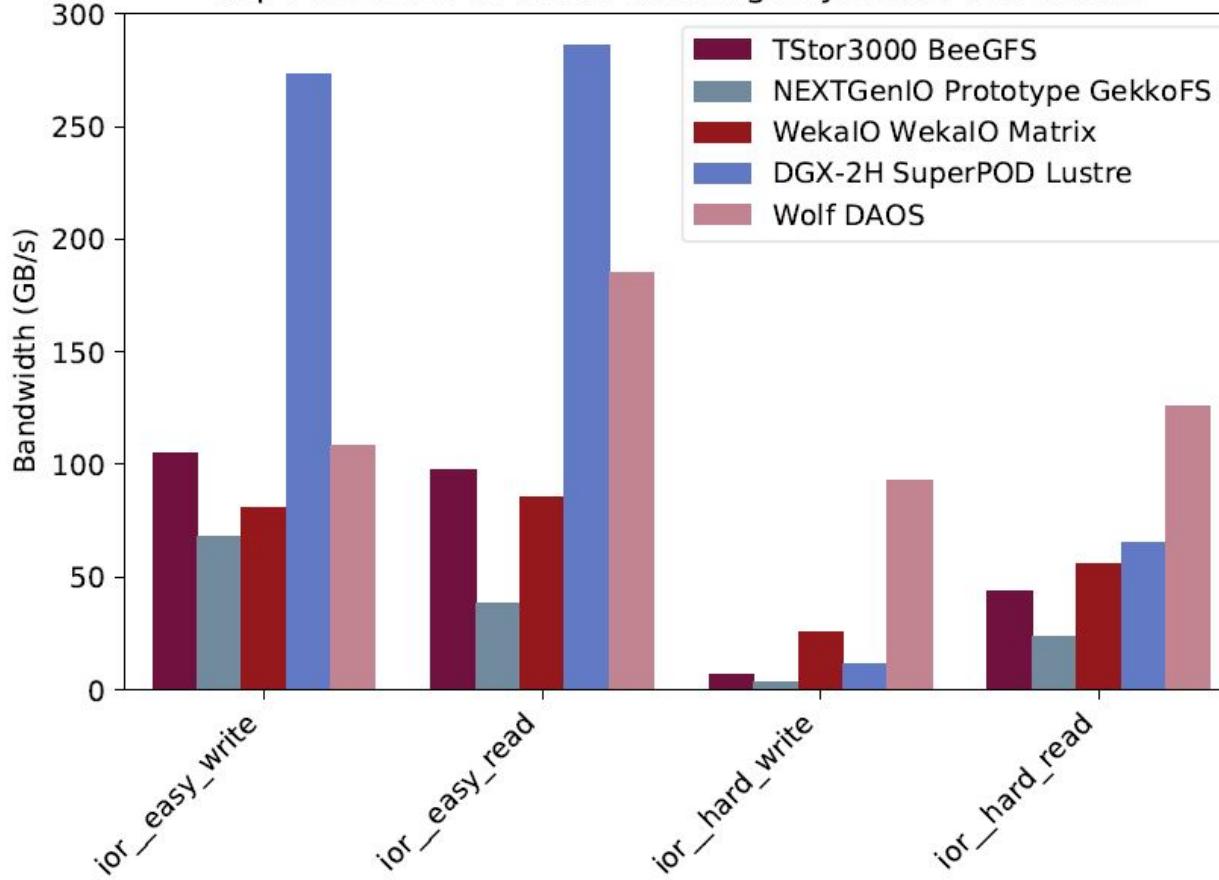
<http://io500.org/list/19-11/>

# SC19 List of Awarded Systems

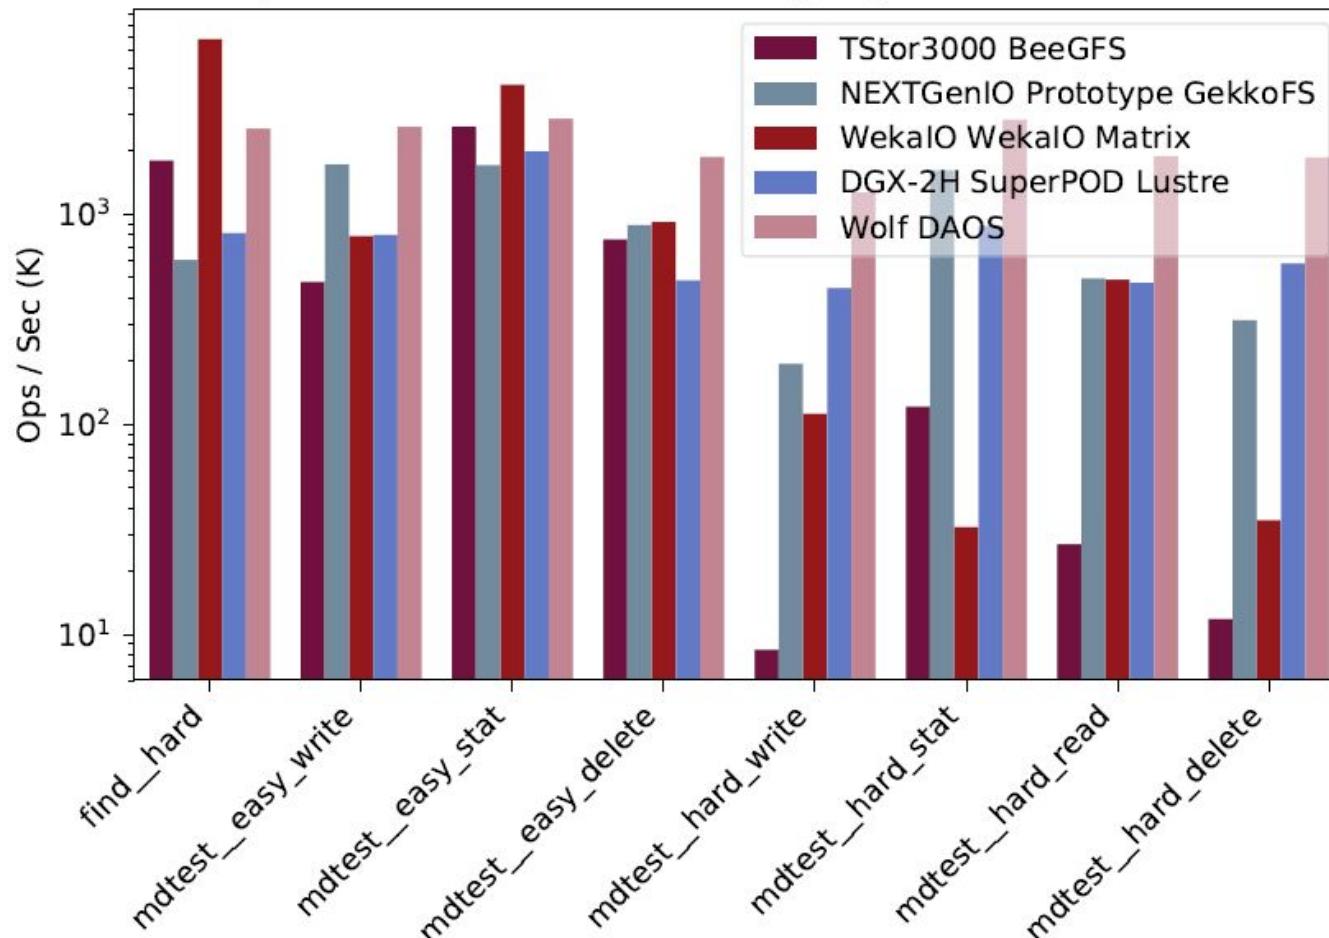
---

10-Node	Metadata	Intel	DAOS	2152	KIOPS
	Bandwidth	Intel	DAOS	124	GiB/s
	Overall	Intel	DAOS	516	score
All Systems	Metadata	WekaIO	Matrix	5045	KIOPS
	Bandwidth	Tianhe	Lustre	209	GiB/s
	Overall	WekaIO	Matrix	939	score

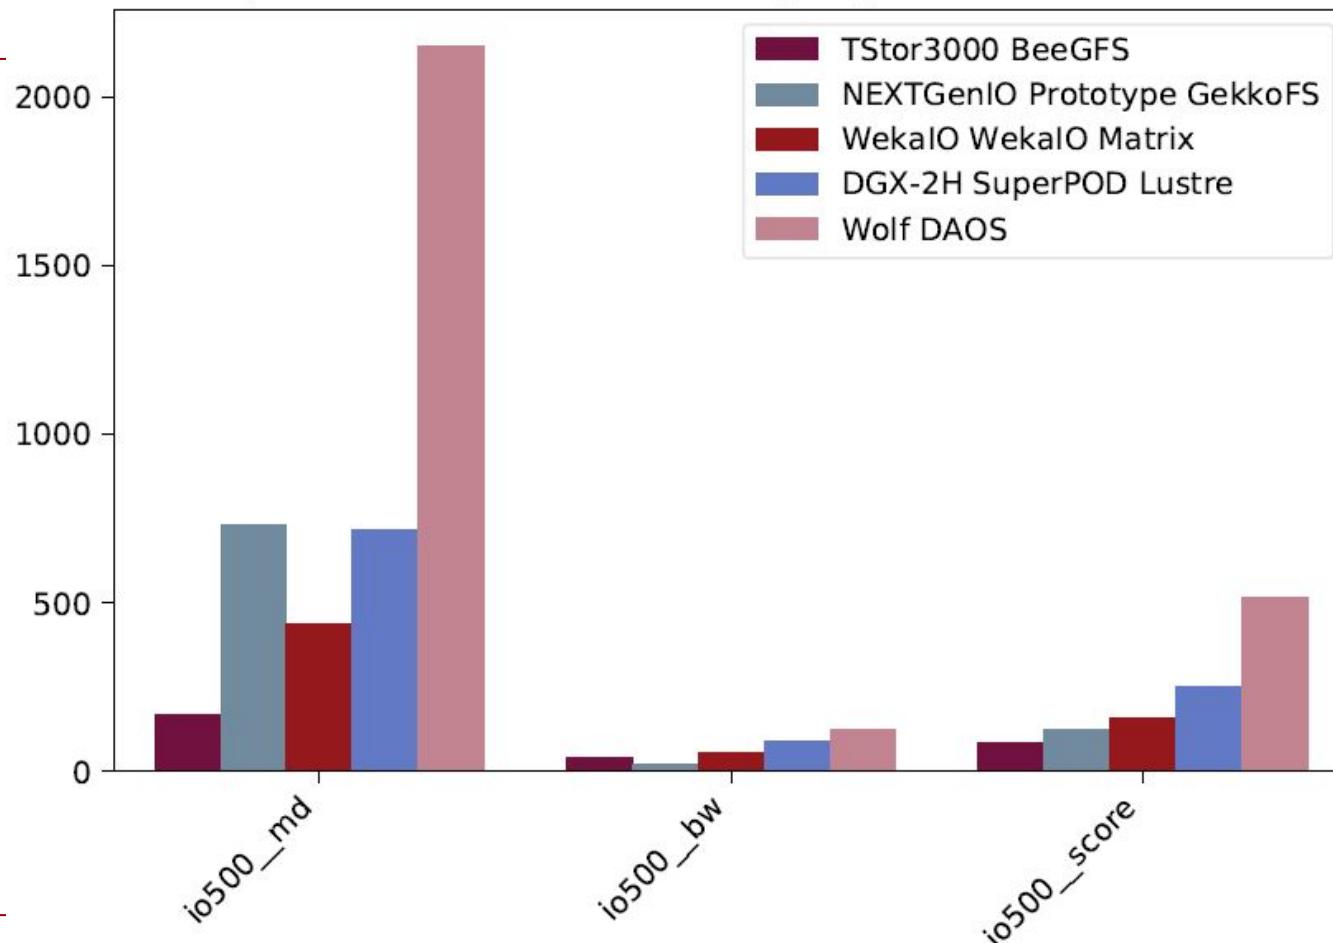
## Top Five SC19 10 Node Challenge Systems: IOR Scores



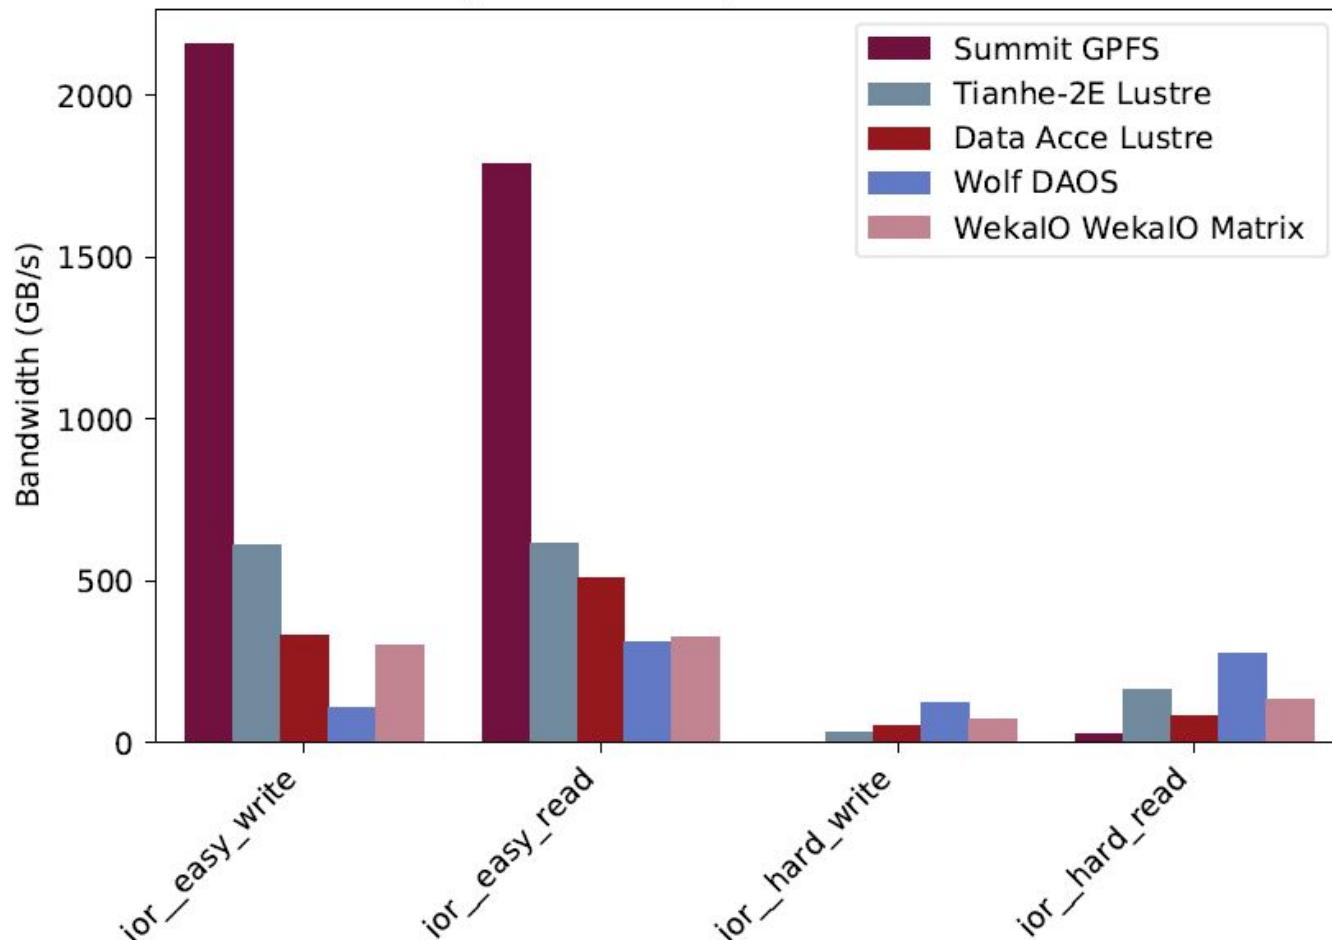
## Top Five SC19 10 Node Challenge Systems: mdtest rates



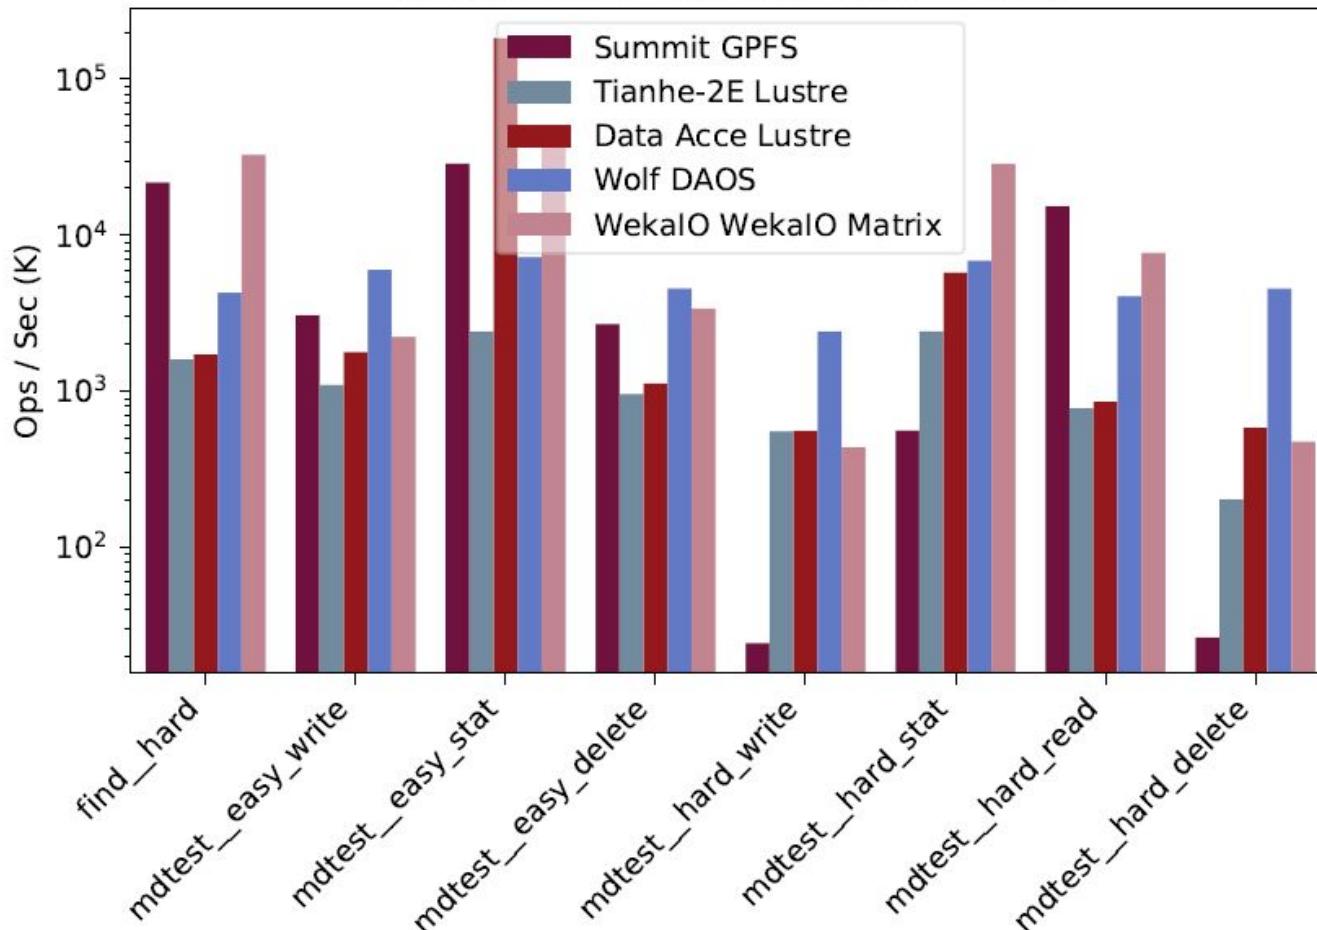
## Top Five SC19 10 Node Challenge Systems: IO500 Scores



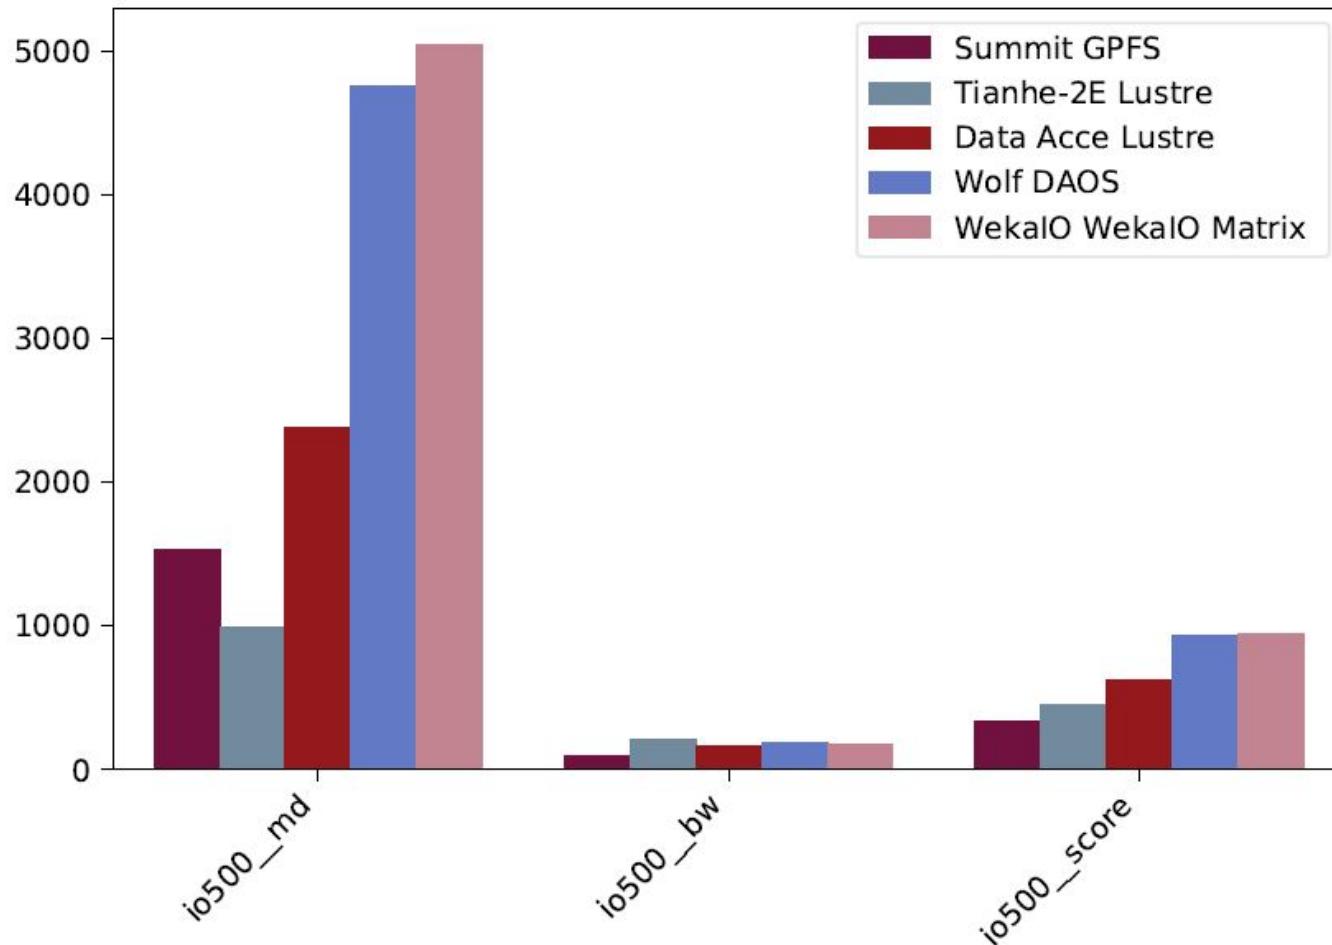
## Top Five SC19 Systems: IOR Scores



## Top Five SC19 Systems: mdtest rates



## Top Five SC19 Systems: IO500 Scores



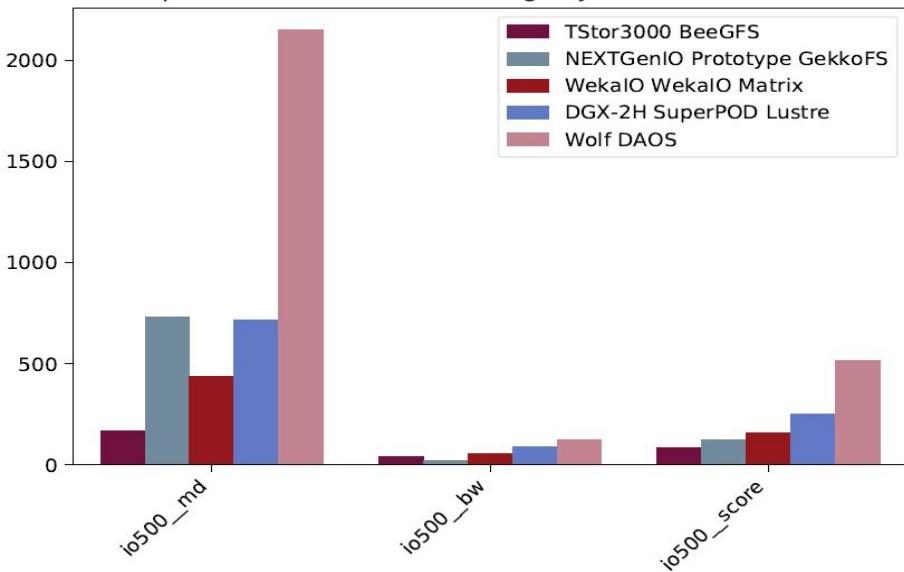
# SC19 IO500

---

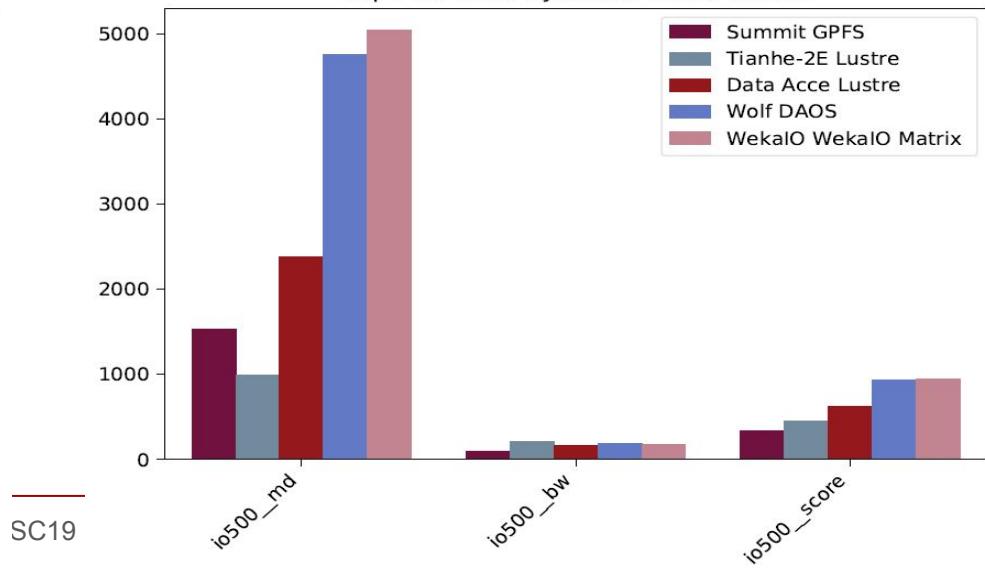
## Winners and Honorable Mentions

10-Node	Metadata	Intel	DAOS	2152	KIOPS
	Bandwidth	Intel	DAOS	124	GiB/s
	<b>Overall</b>	<b>Intel</b>	<b>DAOS</b>	<b>516</b>	<b>score</b>
All Systems	Metadata	WekaIO	Matrix	5045	KIOPS
	Bandwidth	Tianhe	Lustre	209	GiB/s
	<b>Overall</b>	<b>WekaIO</b>	<b>Matrix</b>	<b>939</b>	<b>score</b>

Top Five SC19 10 Node Challenge Systems: IO500 Scores



Top Five SC19 Systems: IO500 Scores



SC19

# SC Student Cluster Competition Preliminary results\*

Ranking	Team	Score	BW	MD	Nodes	DropCache
1	Nanyang	28.28	4.139	193.21	1	yes
2	Purdue	26.0845	9.8	69.4	5	no
3	Shanghai Tech	20.4175	2.61	159.381	1	yes
4	ETH Zurich	19.588	3.23	118.575	1	yes
5	Peking	16.21	2.54	103.2	1	yes
6	FAU	13.82	2.02	94.4	1	yes
7	UIUC	12.69	6.5	24.77	1	yes
8	NTHU	9.4	1.05	84.16	1	yes
9	Tennessee	8.96	1.85	43.38	1	yes
10	Wake Forest	5.44	0.98	30.2	1	yes
11	Warsaw	4.95	1.33	18.38	2	no
12	Washington	4.39	0.248	77.916	1	yes
13	Shanghai Jiao Tong	3.279	1.188	9.05	1	yes

Not valid

	Tsinghua	30.55	3.08	303.121	1	yes
	NC State	24.64	2.39	253.351	1	no

\* Results received less than 18 hours ago and have not been fully validated by the committee yet

# Roadmap

10<sup>500</sup>

# Roadmap for the IO-500

---

- Rewrite the IO-500 into a C-application instead of the script solution
  - Run by using a configuration file and no additional arguments
  - Improved error handling and validation for submitters
  - Produces the same results as the current bash solution
- Using MDTest data validation during mdtest hard read
  - Compares read data with the expectation
- Integration of tools to automatically harvest system configuration
- Rewrite the webpage
  - Move 100% of code into github
  - Mirror at [io500.vi4io.org](http://io500.vi4io.org) and [io500.org](http://io500.org)

# C-Application / Thoughts

---

- Running the application should be as simple as (e.g. SLURM)

```
#!/bin/bash -e
#SBATCH -p compute2
#SBATCH --nodes=10
module load OpenMPI

mpiexec -np 20 ./io500 final-config.ini
```

- Configuration could be INI or JSON files
  - Providing only options that are allowed to tune; options for additional testing in extra section
- The tool provides a dry-run option showing the exact commands it runs
  - e.g. mpiexec -np 20 ./io500 --dry-run final-config.ini
  - Dumps the full INI options that are available and their current values
  - Shows predicted execution behavior:
    - I run mdtest with these arguments, then this then that....
    - The result should be valid or will definitely be invalid based on the options provided

# C-Application / Configuration File Draft

---

```
[find]
external-program = ./bin/mmfnd.sh # wrapper returns similar output

[ior-hard]
API = MPIIO # Like when using ior -O <OPTION>=<VALUE>
hintsFileName = my-hints.txt

[ior-hard write] # Some options might be valid for specific sections
posix.odirect = 1

[optional]
ior-random = enable

[debug] # the program will warn if anything is invalid
drop-caches = TRUE
stonewall-time = 10
```

# Discussion

10<sup>500</sup>

# Edit or add functionalities to IO-500

---

## Change Request

The IO-500 aims to be a robust and long-living benchmark. Nevertheless, the community recognizes the need to consider modifications occasional modifications such as including new access patterns, removing deprecated access patterns, or any other modifications deemed necessary by the community. Therefore, we have established a process to add further benchmarks, which works as follows:

1. A member of the community prepares a (up to) 1-page proposal for the new access pattern to include. This should include a motivation, a rough sketch of the access pattern and justification why the pattern is important. This proposal can then be sent to the community mailing list or the steering board. Deadline: 1 month before the next community meeting – at the moment, these are the birds-of-a-feather sessions at ISC or Supercomputing.
2. The steering board will give feedback to the technical quality of the proposal.
3. The member is given the opportunity to present the proposal at the next following community IO-500 meeting.
4. Given there are no technical concerns, the IO-500 benchmark will be modified for the next submission period to allow the execution of a benchmark that represents the pattern as an \*optional\* benchmarking step. Additionally, the optional field is introduced into subsequent lists and the changes to the benchmark are documented on the webpage.
5. The optional pattern is kept for at least two subsequent IO-500 lists and community meetings. The results and effectiveness of the new pattern are discussed during the community meetings. As a result, it may be removed, remain optional, or may become mandatory.

The committee can be reached at  [committee@io500.org](mailto:committee@io500.org).

<http://io500.org/rules/proposals>

# Open Floor

---

# Issues about Fair Comparisons

---

- Non-erasure vs erasure systems
- Production system versus benchmark-only system
- Vendor submission versus customer submission
- GA File system versus research file system
- Cloud vs on-prem
- Ephemeral vs persistent file system
- Storage media