



The HDF Group



# What NetCDF users should know about HDF5?

Elena Pourmal

The HDF Group

July 20, 2007



# Outline

- The HDF Group and HDF software
- HDF5 Data Model
- Using HDF5 tools to work with NetCDF-4 programs files
- Performance issues
  - ✓ Chunking
  - ✓ Variable-length datatypes
  - ✓ Parallel performance
- Crash proofing in HDF5



# The HDF Group

- Non-for-profit company with a mission to sustain and develop HDF technology affiliated with University of Illinois
- Spun-off NCSA University of Illinois in July 2006
- Located at the U of I Campus South Research Park
- 17 team members, 5 graduate and undergraduate students
- Owns IP for HDF file format and software
- Funded by NASA, DOE, others



# HDF5 file format and I/O library

- General
  - ✓ simple data model
- Flexible
  - ✓ store data of diverse origins, sizes, types
  - ✓ supports complex data structures
- Portable
  - ✓ available for many operating systems and machines
- Scalable
  - ✓ works in high end computing environments
  - ✓ accommodates data of any size or multiplicity
- Efficient
  - ✓ fast access, including parallel i/o
  - ✓ stores big data efficiently



# HDF5 file format and I/O library

- File format
  - ✓ Complex
    - Objects headers
    - Raw data
    - B-trees
    - Local and Global heaps
    - etc
- C Library
  - ✓ 500+ APIs
  - ✓ C++, Fortran90 and Java wrappers
  - ✓ High-level APIs (images, tables, packets)

Apps: simulation, visualization, remote sensing...

Examples: Thermonuclear simulations  
Product modeling  
Data mining tools  
Visualization tools  
Climate models

## Common application-specific data models

appl-specific APIs

NetCDF-4	SAF	hdf5mesh	IDL	HDF-EOS
Unidata/LANL	LLNL, SNL	Grids	COTS	NASA

HDF5 serial & parallel I/O

HDF5 virtual file layer (I/O drivers)

Sec2

Split Files

MPI I/O

Custom

Stream

HDF5 format

Storage



File

Split metadata and raw data files

File on parallel file system

? User-defined device  
Across the network or to/from another application or library



# HDF5 file format and I/O library

---

*For NetCDF-4 users HDF5 complexity is hidden behind  
NetCDF-4 APIs*



# HDF5 Tools

- Command line utilities  
<http://www.hdfgroup.org/hdf5tools.html>
  - Readers
    - ✓ h5dump
    - ✓ h5ls
  - Writers
    - ✓ h5repack
    - ✓ h5copy
    - ✓ h5import
  - Miscellaneous
    - ✓ h5diff, h5repart, h5mkgrp, h5stat, h5debug, h5jam/h5unjam
  - Converters
    - ✓ h52gif, gif2h5, h4toh5, h5toh4
- HDFView (Java browser and editor)



## Other HDF5 Tools

- ✓ HDF Explorer

Windows only, works with NetCDF-4 files

<http://www.space-research.org/>

- ✓ PyTables

- ✓ IDL

- ✓ Matlab

- ✓ Labview

- ✓ Mathematica

- ✓ See

<http://www.hdfgroup.org/tools5app.html>



# HDF Information

- HDF Information Center  
<http://hdfgroup.org>
- HDF Help email address  
[help@hdfgroup.org](mailto:help@hdfgroup.org)
- HDF users mailing lists  
[news@hdfgroup.org](mailto:news@hdfgroup.org)  
[hdf-forum@hdfgroup.org](mailto:hdf-forum@hdfgroup.org)



# NetCDF and HDF5 terminology

NetCDF	HDF5
Dataset	HDF5 file
Dimensions	Dataspace
Attribute	Attribute
Variable	Dataset
Coordinate variable	Dimension scale



HDFView

File Window Tools Help

File/URL C:\data\HDF projects\ENSIGHT\converter\h5\_EnSight package\hdf5\_ensight\hdf5\_files\ami.h5

ami.h5

EnSight\_model

- geometry
  - Open
  - Open As
  - New
- Copy
- Paste
- Delete
- Save to
- Rename
- Show Properties
- Show Properties
- Close File

AMI-X Blended Body -- 25 X 17 X 49 mesh

Properties - /EnSight\_model/geometry

General Attributes

Number of attributes = 2

Name	Value	Type	Array Size
description_1	AMI-X Blended B...	String, length = 39	1
description_2	Mach .8, 8 degree...	String, length = ...	1

Close

0825x3, stride1x1 ]

Log Info Metadata



The HDF Group



# HDF5 Data Model



# HDF5 data model

- HDF5 file – container for scientific data
- Primary Objects
  - Groups
  - Datasets
- Additional ways to organize data
  - Attributes
  - Sharable objects
  - Storage and access properties

## Metadata

Dataspace	
Rank	Dimensions
3	Dim_1 = 4
	Dim_2 = 5
	Dim_3 = 7

## Datatype

IEEE 32-bit float

## Storage info

chunked

compressed

checksum


## Attributes

time = 32.4

pressure = 987


temp = 56

## Data




- HDF5 atomic types
  - ✓ normal integer & float
  - ✓ user-definable (e.g. 13-bit integer)
  - ✓ variable length types (e.g. strings, ragged arrays)
  - ✓ pointers - references to objects/dataset regions
  - ✓ enumeration - names mapped to integers
  - ✓ array
  - ✓ opaque
- HDF5 compound types
  - ✓ Comparable to C structs
  - ✓ Members can be atomic or compound types
  - ✓ No restriction on complexity

# HDF5 dataset: array of records



Dimensionality:  $5 \times 3$


Datatype:



Record

# Groups

- A mechanism for collections of related objects
- Every file starts with a root group
- Similar to UNIX directories
- Can have attributes
- Objects are identified by a path e.g. /d/b, /t/a





# Attributes

- Attribute – data of the form “name = value”, attached to an object (group, dataset, named datatype)
- Operations scaled down versions of dataset operations
  - ✓ Not extendible
  - ✓ No compression
  - ✓ No partial I/O
- Optional
- Can be overwritten, deleted, added during the “life” of a dataset
- Size under 64K in releases before HDF5 1.8.0



The HDF Group



# Using HDF5 tools with NetCDF-4 programs and files



# Example

- Create netCDF-4 file
- /Users/epourmal/Working/\_NetCDF-4
  - s.c creates simple\_xy.nc (NetCDF3 file)
  - sh5.c creates simple\_xy\_h5.nc (NetCDF4 file)
  - Use h5cc script to compile both examples
  - See contents simple\_xy\_h5.nc with ncdump and h5dump
    - Useful flags
      - ✓ -h to print help menu
      - ✓ -b to export data to binary file
      - ✓ -H to display metadata information only
- HDF Explorer



# NetCDF view: ncdump output

```
% ncdump -h simple_xy_h5.nc
```

```
netcdf simple_xy_h5 {
```

```
dimensions:
```

```
    x = 6 ;
```

```
    y = 12 ;
```

```
variables:
```

```
    int data(x, y) ;
```

```
data:
```

```
}
```

```
% h5dump -H simple_xy.nc
```

```
h5dump error: unable to open file "simple_xy.nc"
```

- ✓ This is NetCDF3 file, h5dump will not work




# HDF5 view: h5dump output

```
% h5dump -H simple_xy_h5.nc
HDF5 "simple_xy_h5.nc" {
GROUP "/" {
DATASET "data" {
DATATYPE H5T_STD_I32LE
DATASPACE SIMPLE {(6,12) / (6,12)}
ATTRIBUTE "DIMENSION_LIST" {
DATATYPE H5T_VLEN { H5T_REFERENCE}
DATASPACE SIMPLE {(2) / (2)}
}
}
DATASET "x" {
DATATYPE H5T_IEEE_F32BE
DATASPACE SIMPLE {(6) / (6)}
.....
}
```



# HDF Explorer





# HDF Explorer

**HDF Explorer - [simple\_xy\_h5.nc:data]**

File Edit View Options Window Help

simple\_xy\_h5.nc

data

DIMENSION\_LIST

x

CLASS  
NAME  
REFERENCE\_L

y

CLASS  
NAME  
REFERENCE\_L

	0	1	2	3	4	5	6	7	8	9
0	0	1	2	3	4	5	6	7	8	9
1	12	13	14	15	16	17	18	19	20	21
2	24	25	26	27	28	29	30	31	32	33
3	36	37	38	39	40	41	42	43	44	45
4	48	49	50	51	52	53	54	55	56	57
5	60	61	62	63	64	65	66	67	68	69

Ready Layer 1 of 1



The HDF Group




# Performance issues



# Performance issues

- Choose appropriate HDF5 library features to organize and access data in HDF5 files
- Three examples:
  - Collective vs. Independent access in parallel HDF5 library
  - Chunking
  - Variable length data

# Layers – parallel example





# h5perf

- An I/O performance measurement tool
- Test 3 File I/O API
  - Posix I/O (open/write/read/close...)
  - MPIO (MPI\_File\_{open,write,read.close})
  - PHDF5
    - H5Pset\_fapl\_mpio (using MPI-IO)
    - H5Pset\_fapl\_mpio\_posix (using Posix I/O)



## H5perf: Some features


- Check (-c) verify data correctness
- Added 2-D chunk patterns in v1.8




## My PHDF5 Application I/O “inhales”

- If my application I/O performance is bad, what can I do?
  - Use larger I/O data sizes
  - Independent vs Collective I/O
  - Specific I/O system hints
  - Parallel File System limits

- User reported Independent data transfer was much slower than the Collective mode
- Data array was tall and thin:  
230,000 rows by 6 columns



230,000 rows






# Independent vs. Collective write

(6 processes, IBM p-690, AIX, GPFS)

# of Rows	Data Size (MB)	Independent (Sec.)	Collective (Sec.)
16384	0.25	8.26	1.72
32768	0.50	65.12	1.80
65536	1.00	108.20	2.68
122918	1.88	276.57	3.11
150000	2.29	528.15	3.63
180300	2.75	881.39	4.12


# Independent vs Collective write (6 processes, IBM p-690, AIX, GPFS)





# Some performance results

1. A parallel version of NetCDF-3 from ANL/Northwestern University/University of Chicago (PnetCDF)
2. HDF5 parallel library 1.6.5
3. NetCDF-4 beta1
4. For more details see  
<http://www.hdfgroup.uiuc.edu/papers/papers/ParallelPerformance.pdf>




Flash I/O Website [http://flash.uchicago.edu/~zingale/flash\\_benchmark\\_io/](http://flash.uchicago.edu/~zingale/flash_benchmark_io/)


Robb Ross, etc."Parallel NetCDF: A Scientific High-Performance I/O Interface




# HDF5 and PnetCDF performance comparison




Bluesky: Power 4



uP: Power 5




Bluesky: Power 4




uP: Power 5

# Parallel NetCDF-4 and PnetCDF




- Fixed problem size = 995 MB
- Performance of PnetCDF4 is close to PnetCDF

# HDF5 chunked dataset




- Dataset is partitioned into fixed-size chunks
- Data can be added along any dimension
- Compression is applied to each chunk
- Datatype conversion is applied to each chunk
- *Chunking storage creates additional overhead in a file*
- *Do not use small chunks*


# Writing chunked dataset



# Writing chunked datasets





# Partial I/O for chunked dataset



- Build list of chunks and loop through the list
  - For each chunk:
    - Bring chunk into memory
    - Map selection in memory to selection in file
    - Gather elements into conversion buffer and perform conversion
    - Scatter elements back to the chunk
    - Apply filters (compression) when chunk is flushed from chunk cache
- For each element 3 memcpy performed*

# Partial I/O for chunked dataset






On eviction from cache chunk is compressed  
and is written to the file




Great performance



Selection coincides with a chunk

Poor performance



Selection spans over all chunks



## Things to remember about HDF5 chunking

- ✓ Use appropriate chunk sizes
- ✓ Make sure that cache is big enough to contain chunks for partial I/O
- ✓ Use hyperslab selections that are aligned with chunks
- ✓ Memory may grow when application opens and modifies a lot of chunked datasets



## Variable length datasets and I/O

- Examples of variable-length data

- String

A[0] “*the first string we want to write*”

.....

A[N-1] “*the N-th string we want to write*”

- Each element is a record of variable-length

A[0] (1,1,0,0,0,5,6,7,8,9) length of the first record is 10

A[1] (0,0,110,2005)

.....

A[N] (1,2,3,4,5,6,7,8,9,10,11,12,...,M) length of the N+1 record is M


- Variable length description in HDF5 application

```
typedef struct {  
    size_t length;  
    void *p;  
} hvl_t;
```


- Base type can be any HDF5 type

```
H5Tvlen_create(base_type)
```

- ~ 20 bytes overhead for each element
- Raw data cannot be compressed



## VL chunked dataset in a file





# Variable length datasets and I/O

- Hints
  - Avoid closing/opening a file while writing VL datasets
    - global heap information is lost
    - global heaps may have unused space
  - Avoid writing VL datasets interchangeably
    - data from different datasets will be written to the same heap
  - If maximum length of the record is known, use fixed-length records and compression



The HDF Group



# Crash-proofing



# Why crash proofing?

- HDF5 applications tend to run long times (sometimes until system crashes)
- Application crash may leave HDF5 file in a corrupted state
- Currently there is no way to recover data
- One of the main obstacles for production codes that use NetCDF-3 to move to NetCDF-4
- Funded by ASC project
- Prototype release is scheduled for the end of 2007



# HDF5 Solution

- Journaling
  - Modifications to HDF5 metadata are stored in an external journal file
  - HDF5 will be using asynchronous writes to the journal file for efficiency
- Recovering after crash
  - HDF5 recovery tool will replay the journal and apply all metadata writes bringing HDF5 file to a consistent state
  - Raw data will consist of data that made to disk
- Solution will be applicable for both sequential and parallel modes



Thank you!

Questions ?