


# Scientific I/O: more than hard drives and file systems

- Scientists think about data in terms of their science problem: molecules, atoms, grid cells, particles
- Ultimately, physical disks store bytes of data
- Layers in between, the application and physical disks are at various levels of sophistication



Images from David Randall, Paola Cessi, John Bell, T Scheibe



# Example: FLASH Astrophysics

- FLASH is an astrophysics code for studying events such as supernovae

- Adaptive-mesh hydrodynamics
- Scales to 1000s of processors
- MPI for communication

- Frequently checkpoints:

- Large blocks of typed variables from all processes
- Portable format
- Canonical ordering (different than in memory)
- Skipping ghost cells



# Example: FLASH's HDF5 requirements

- FLASH AMR structures do not map directly to HDF5 multidimensional arrays
- Must create mapping of the in-memory FLASH data structures into a representation in HDF5 multidimensional arrays
- Chose to
  - Place all checkpoint data in a single file
  - Impose a linear ordering on the AMR blocks
    - Use 4D variables
  - Store each FLASH variable in its own HDF5 variable
    - Skip ghost cells
  - Record attributes describing run time, total blocks, etc.

# FLASH HDF5 Usage

## ■ Annotations describing data, experiment

```
attribute_id = H5Acreate(group_id,  
    "iteration",H5T_NATIVE_INT, dataspace_id, H5P_DEFAULT);  
status = H5Awrite(attribute_id, H5T_NATIVE_INT, temp);
```

## ■ HDF5 variables for each FLASH variable

```
ierr = H5Sselect_hyperslab(dataspace, H5S_SELECT_SET,  
    start_4d, stride_4d, count_4d, NULL);  
status = H5Dwrite(dataset, H5T_NATIVE_DOUBLE, memspace,  
    dataspace, dxfer_template, unknowns);
```

# HACC: understanding cosmos via simulation

- “Cosmology = Physics + Simulation” (Salman Habib)
- Sky surveys collecting massive amounts of data
  - (~100 PB)
- Understanding of these massive datasets rests on modeling distribution of cosmic entities
- Seed simulations with initial conditions
- Run for 13 billion (simulated) years
- Comparison with observed data validates physics model.
- I/O challenges:
  - Checkpointing
  - analysis

