



BERKELEY LAB
Bringing Science Solutions to the World



**U.S. DEPARTMENT OF
ENERGY**



Autocorrelation with ADIOS



What is ADIOS

An extendable **framework** that allows developers to *plug-in*

- **I/O methods:** N-to-M, N-to-N, N-to-1, In Situ (aka Staging)
- **Transformations:** Compression, Decompression, Indexing
- **Self describing** data format: ADIOS-BP
- **Indexing/Querying:** MinMax, FastBit, Alacrity

Incorporates the “best” practices in the I/O middleware layer

Released twice a year, now 1.12, under the completely free BSD license

- <https://www.olcf.ornl.gov/center-projects/adios>
- <https://github.com/ornladios/ADIOS>

Available at ALCF, OLCF, NERSC, CSCS, Tianhe-1,2, Pawsey SC, Ostrava

Applications are supported through OLCF INCITE program

Outreach via on-line manuals, and live tutorials

How to use ADIOS

ADIOS is provided as a library to users; use it like other I/O libraries, except

ADIOS has a **simple approach** for I/O

- User defines in application source code: “**what**” and “**when**”
 - **Every process defines what data and when to output**
- ADIOS takes care of the “**how**”

Biggest hurdle for users:

- Forget all of your manual **tricks** to gain I/O performance on your particular target system and target scale and just say what you want to write/read
- Trust ADIOS to deliver the performance

Performance Portability:

- Write once, **perform well anywhere**
 - It comes naturally with ADIOS
 - ADIOS has many different I/O methods (strategies)
-

Data management tradeoffs at exascale → to hybrid staging

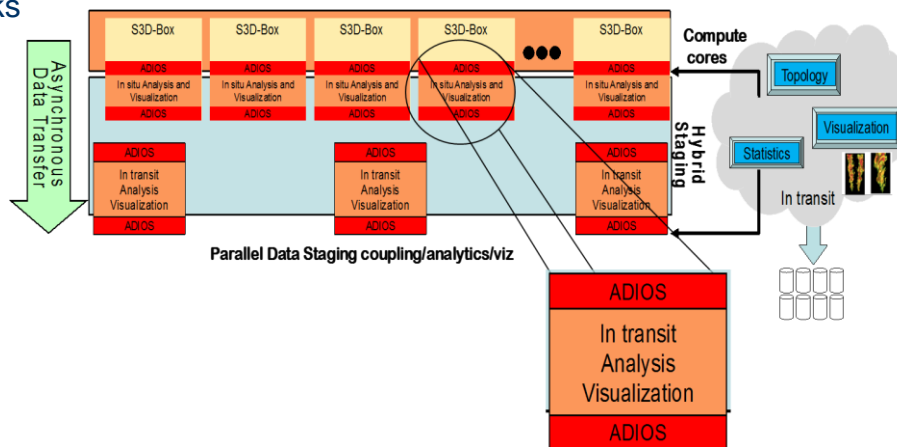
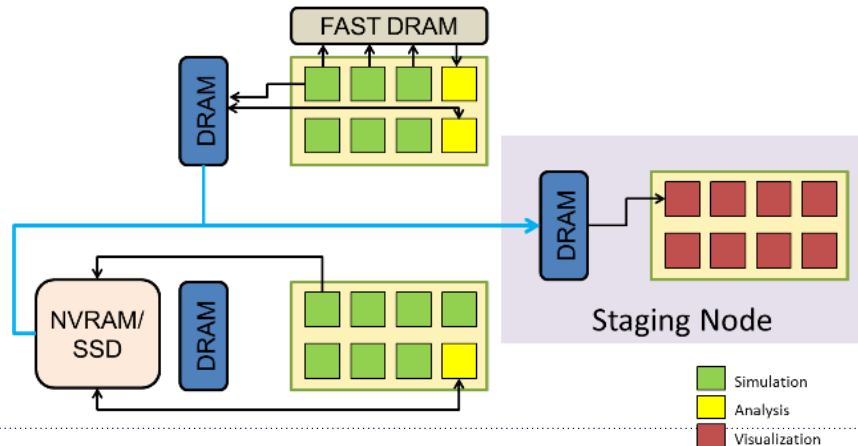
Explore node layout choices for data management

Balance of memory size and speed

Feedback for node designs with NVRAM, larger memory, on-chip NIC

Network throughput and latency impact on SDMA tasks

Placement of operations in concert with solver and network topology



Goals of the ADIOS Read API design

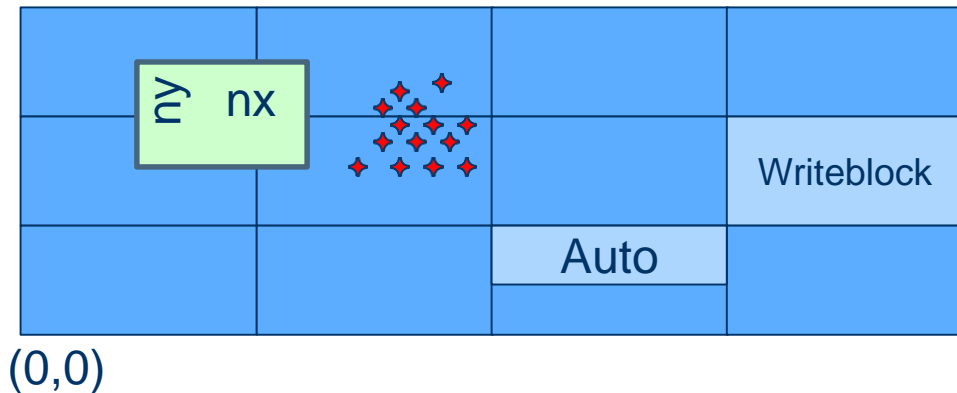
Staging I/O

- Insulate the scalable application from the **variability** inherent in the file system
- Enable the utilization of **in situ and in-transit analytics and visualization**

Same API for reading data from files and from staging

Allow for read optimizations:

- **Multiple read** operations can be scheduled before performing them
- Allow for blocking and **non-blocking** reads
- Use generic **selections** in the read statements instead of describing a bounding box
- Option to let ADIOS deliver data in **chunks**, with memory allocated inside ADIOS not in user-space



Selections

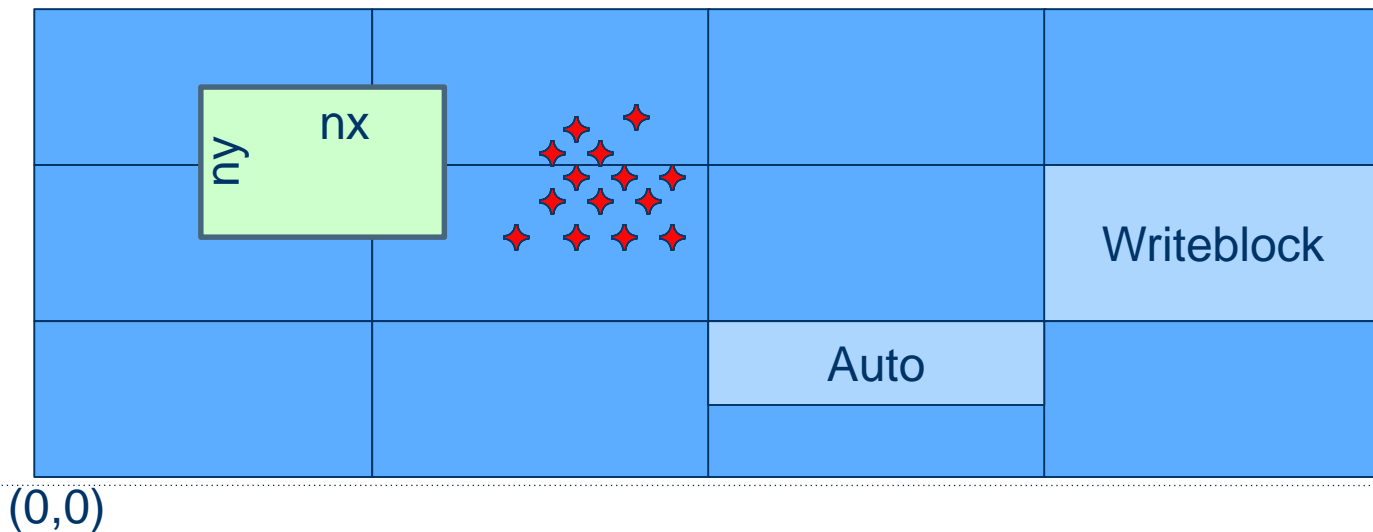
ADIOS_SELECTION *

adios_selection_boundingBox (int ndim, uint64_t * offsets, uint64_t * readsize)

adios_selection_points (uint64_t ndim, uint64_t npoints, uint64_t *points)

adios_selection_writeblock (int index)

adios_selection_auto (char * hints)



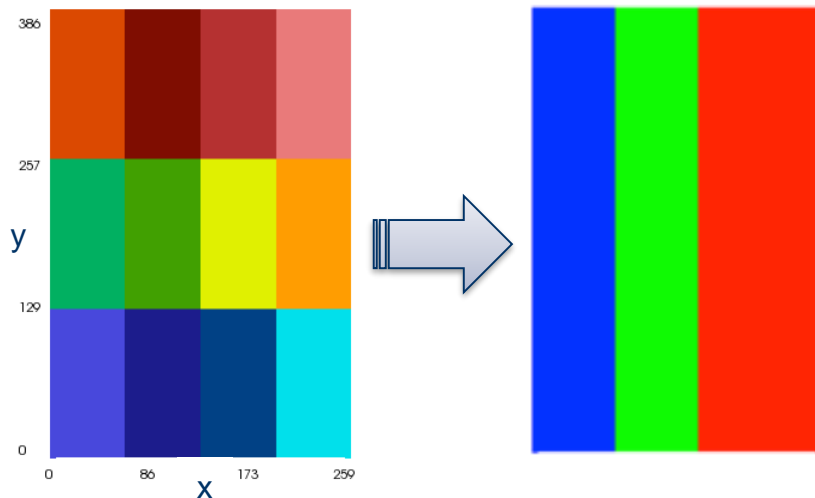
Example of Read API: read a variable step-by-step

```
int count[] = {10,10,10};  
int offs[] = {5,5,5};  
  
P = (double*) malloc (sizeof(double) * count[0] * count[1] * count[2]);  
Q = (double*) malloc (sizeof(double) * count[0] * count[1] * count[2]);  
ADIOS_SELECTION *sel = adios_select_boundingBox (3, offs, count);  
while (fp != NULL) {  
    adios_schedule_read (fp, sel, "P", 0, 1, P);  
    adios_schedule_read (fp, sel, "Q", 0, 1, Q);  
    adios_perform_reads (fp, 1, NULL); // 1: blocking read  
  
    // P and Q contains the data at this point  
    adios_release_step (fp); // staging method can release this step  
  
    // ... process P and Q, then advance the step  
    adios_advance_step (fp, 0, 60.0);  
  
    // 60 sec blocking wait for the next available step  
}  
  
// free ADIOS resources  
adios_free_selection (sel);
```

N to M reorganization with stage_write

heat transfer + stage_write running together

- Write out 6 time-steps.
- Write from 12 cores, arranged in a 4 x 3 arrangement.
- Read from 3 cores, arranged as 1x3



N to M reorganization with stage_write

```
$ cd ~/Tutorial/heat_transfer
edit heat_transfer.xml (vi, gedit)
set method to MPI
❏ <method group="heat" method="MPI"/>
```

```
$ mpirun -np 12 ./heat_transfer_adios1 heat 4 3 40 50 6 500
```

```
$ bpls -D heat.bp T
```

```
double T 6*{150, 160}
step 0:
  block 0: [ 0: 49, 0: 39]
  block 1: [ 0: 49, 40: 79]
  ...
  block 11: [100:149, 120:159]
```

```
$ mpirun -np 3 stage_write/stage_write heat.bp h_3.bp BP "" FLEXPATH "" 3
```

```
$ bpls -D h_3.bp T
```

```
double T 6*{150, 160}
step 0:
  block 0: [ 0:149, 0: 52]
  block 1: [ 0:149, 53:105]
  block 2: [ 0:149, 106:159]
```

Live demo

- Live demo on virtual machine